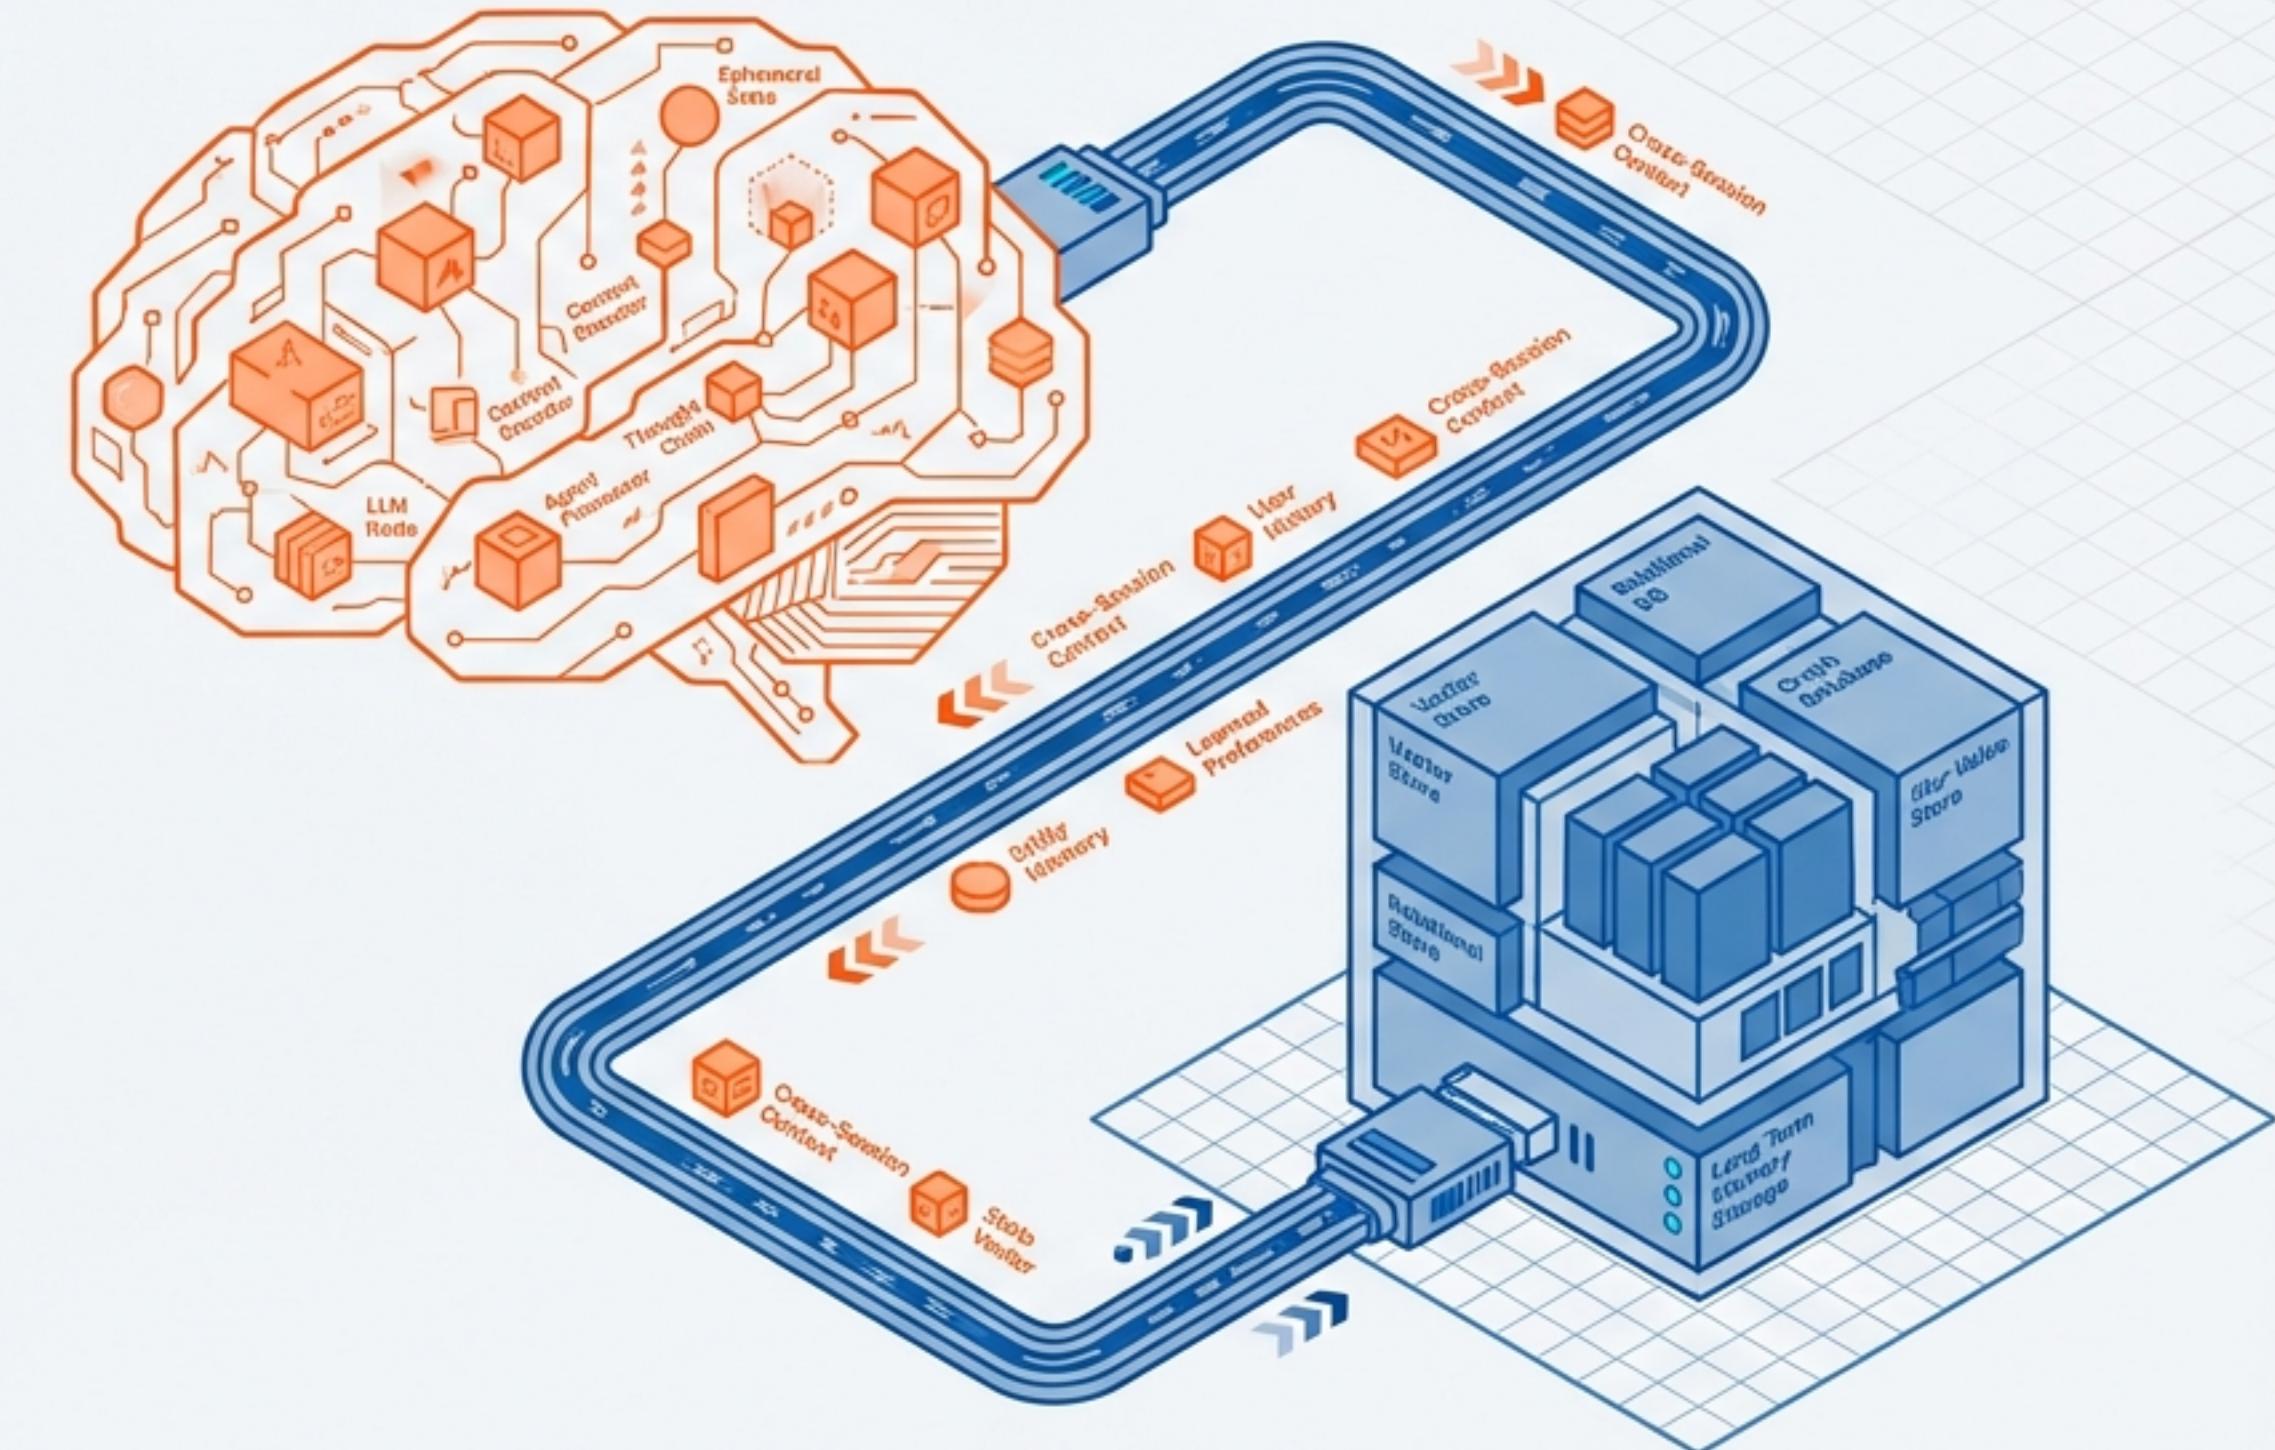


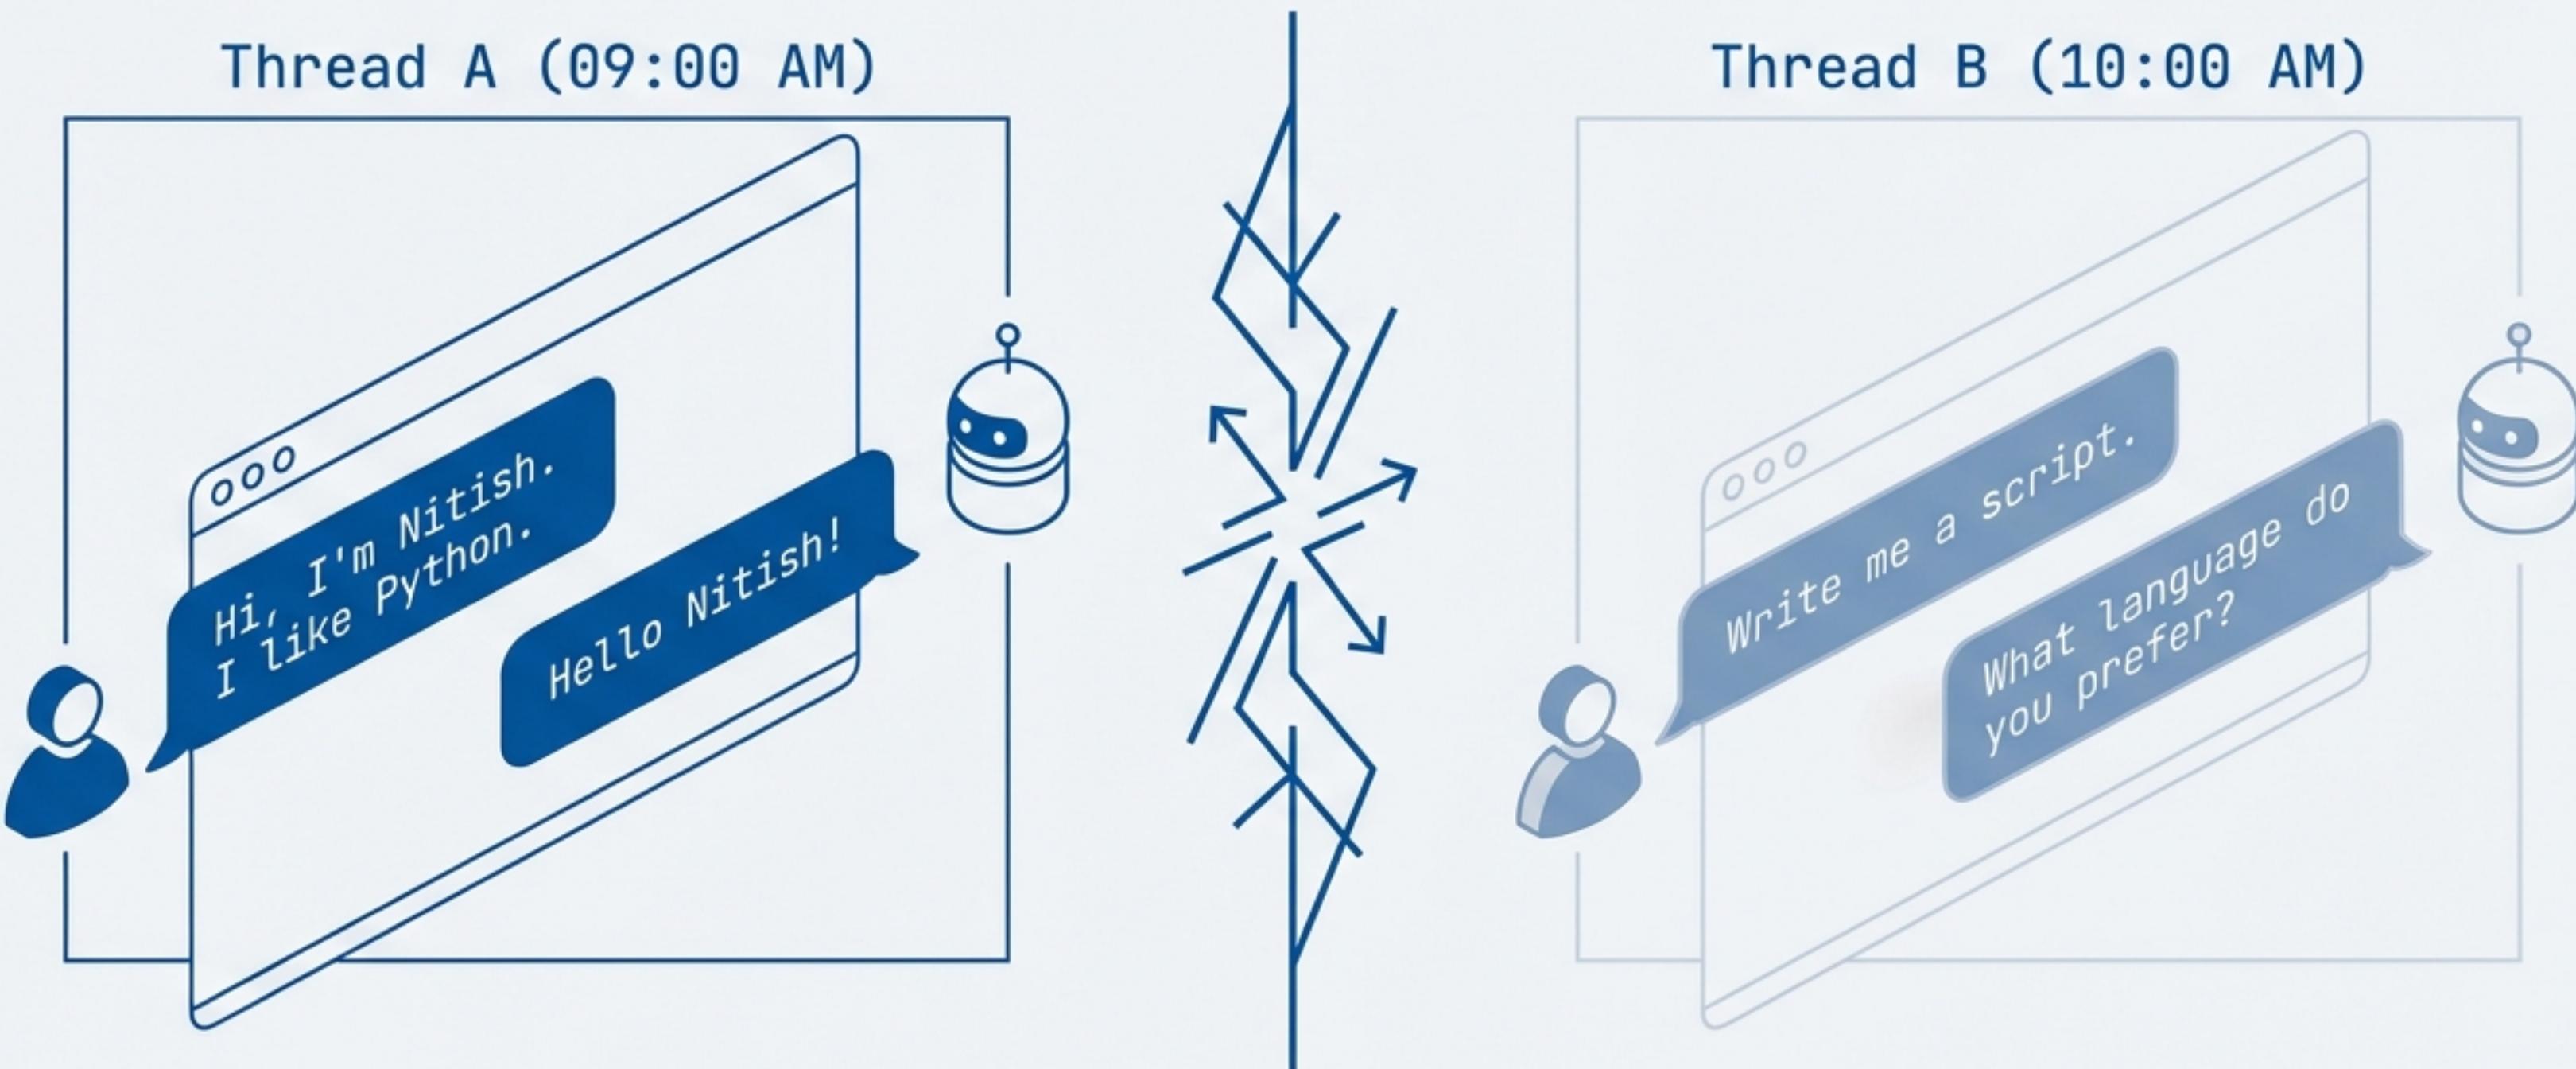
# Beyond the Thread: Architecting Long Term Memory

Implementing persistent,  
cross-session context in  
LangGraph agents.



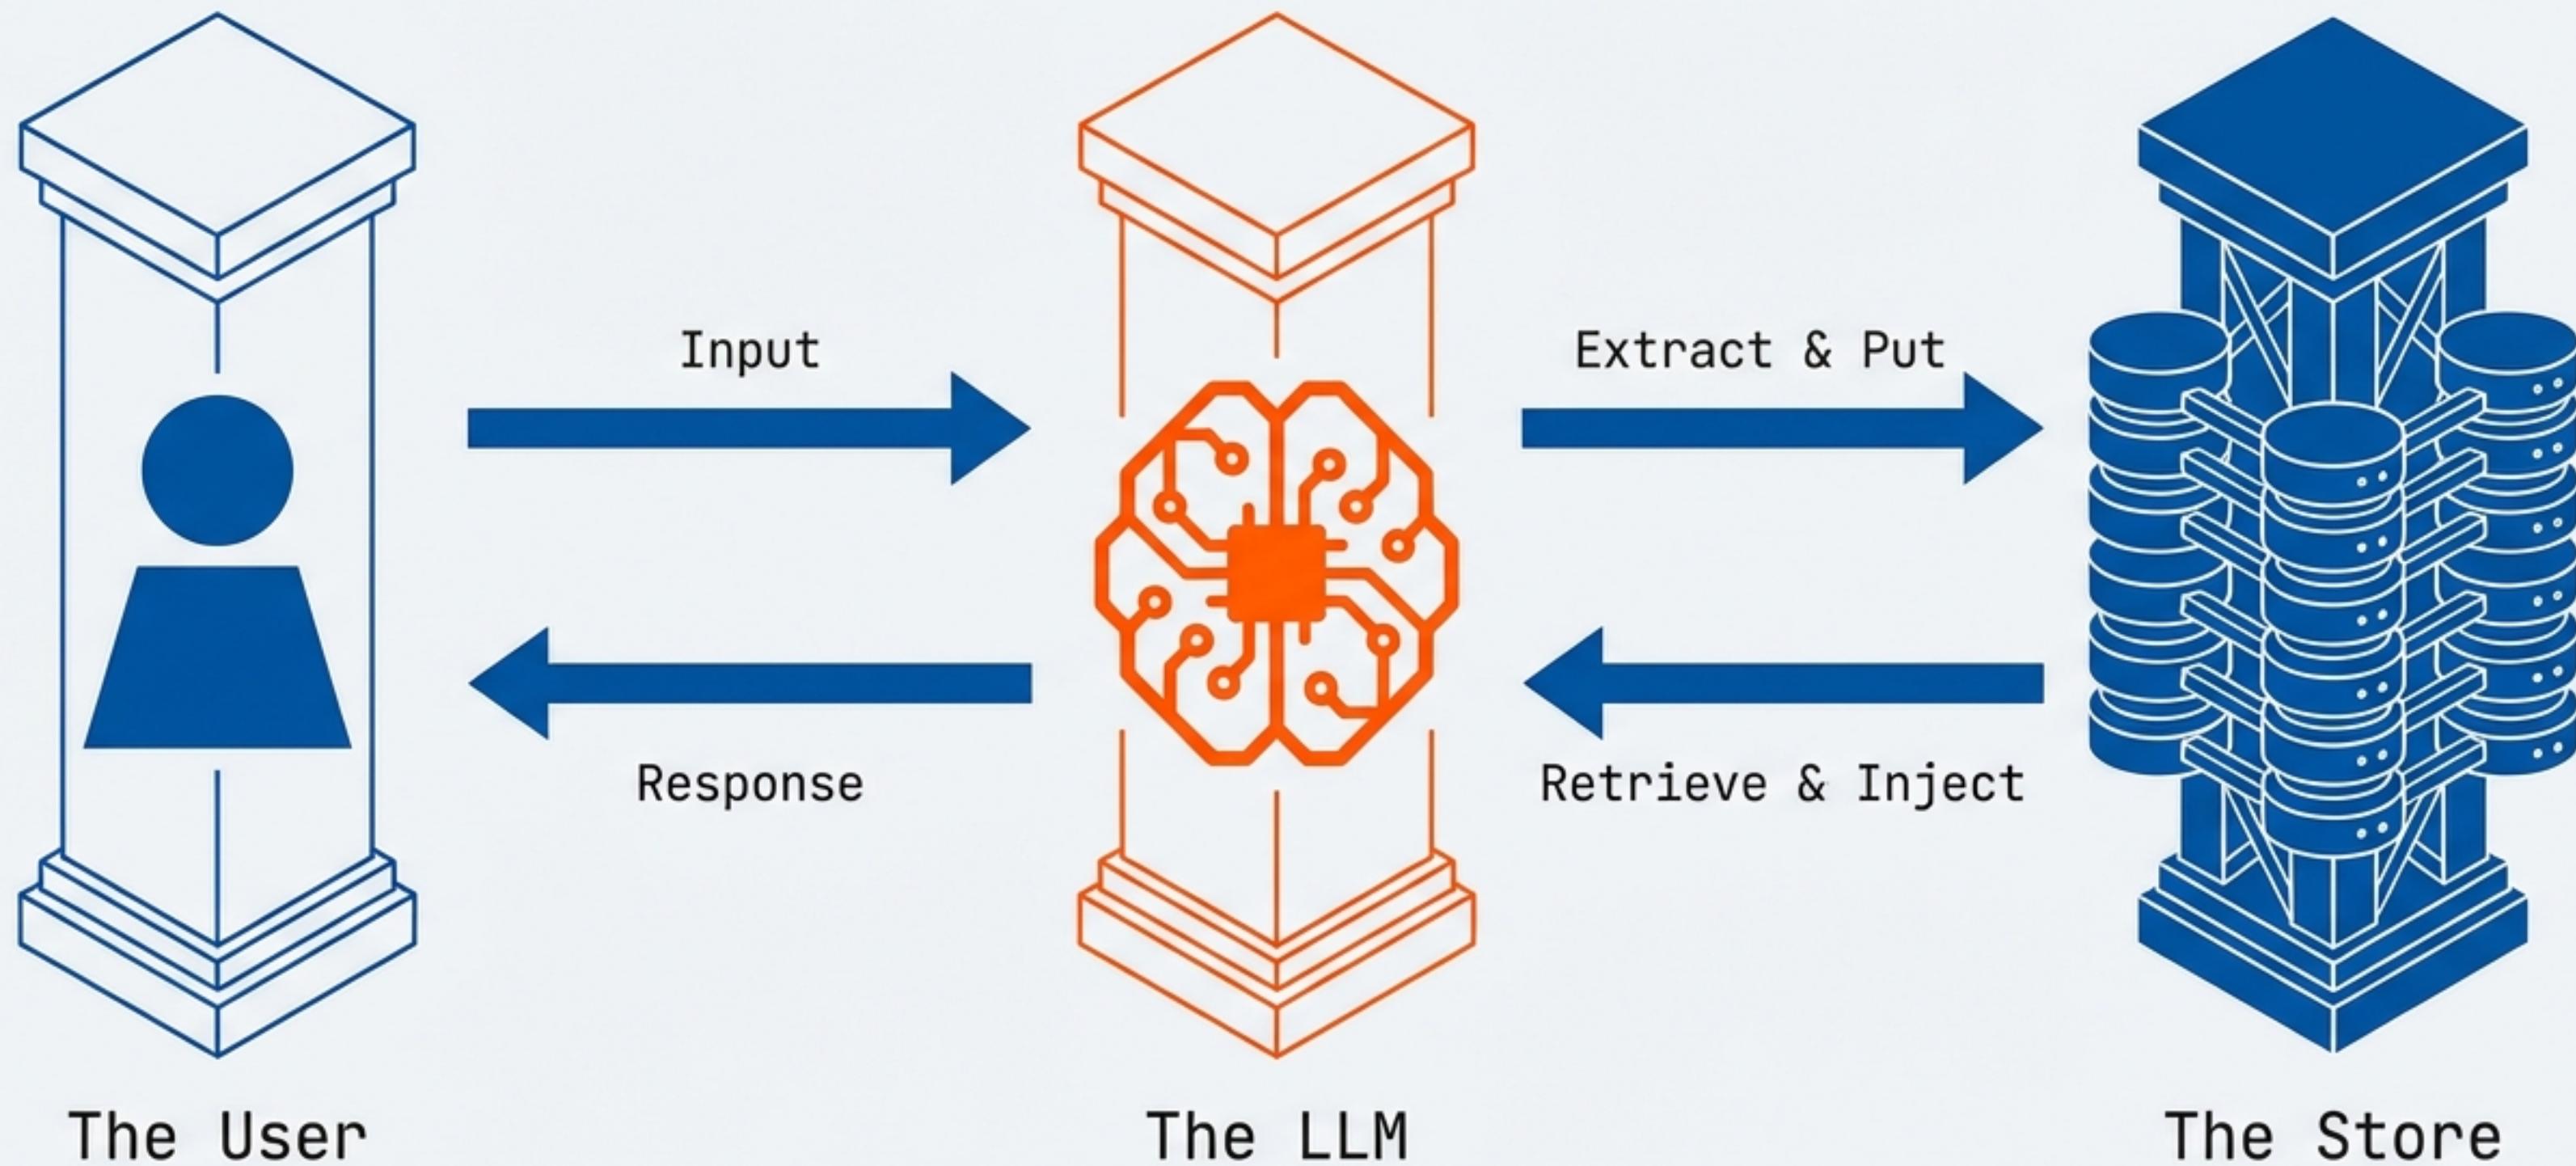
Based on the technical implementation by CampusX

# The Problem of Ephemeral Context



**LLM context is thread-scoped. Closing the window resets the relationship.**

# The Third Pillar of Architecture

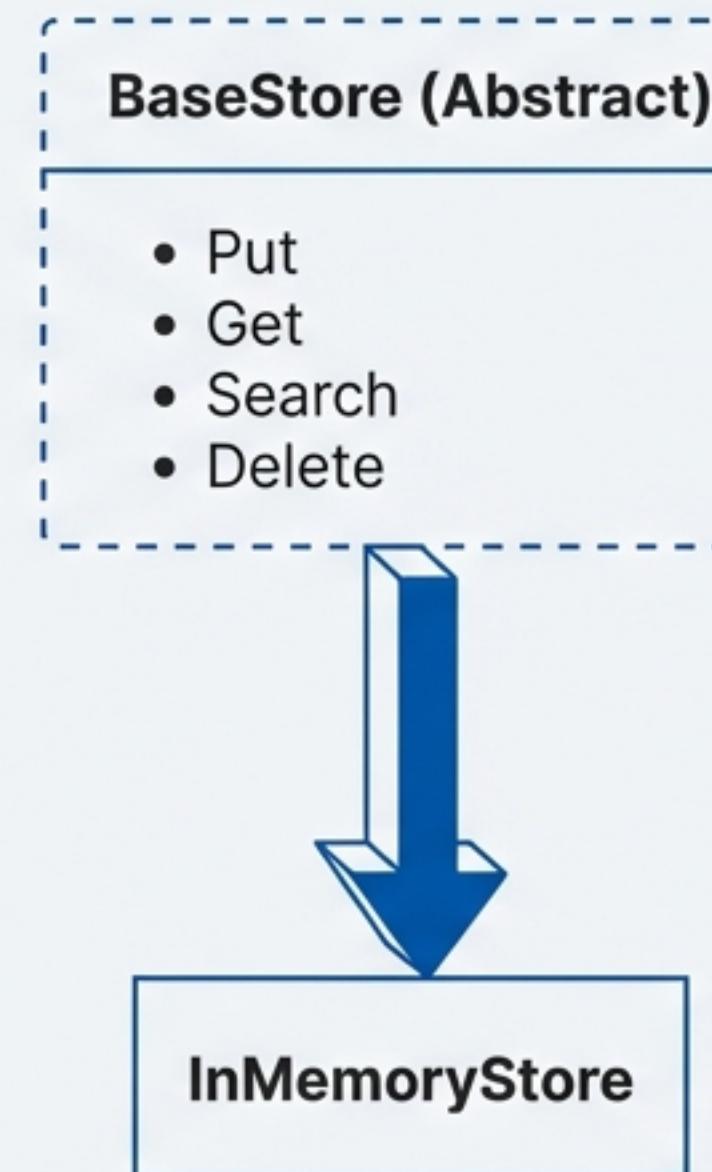


The User

The LLM

The Store

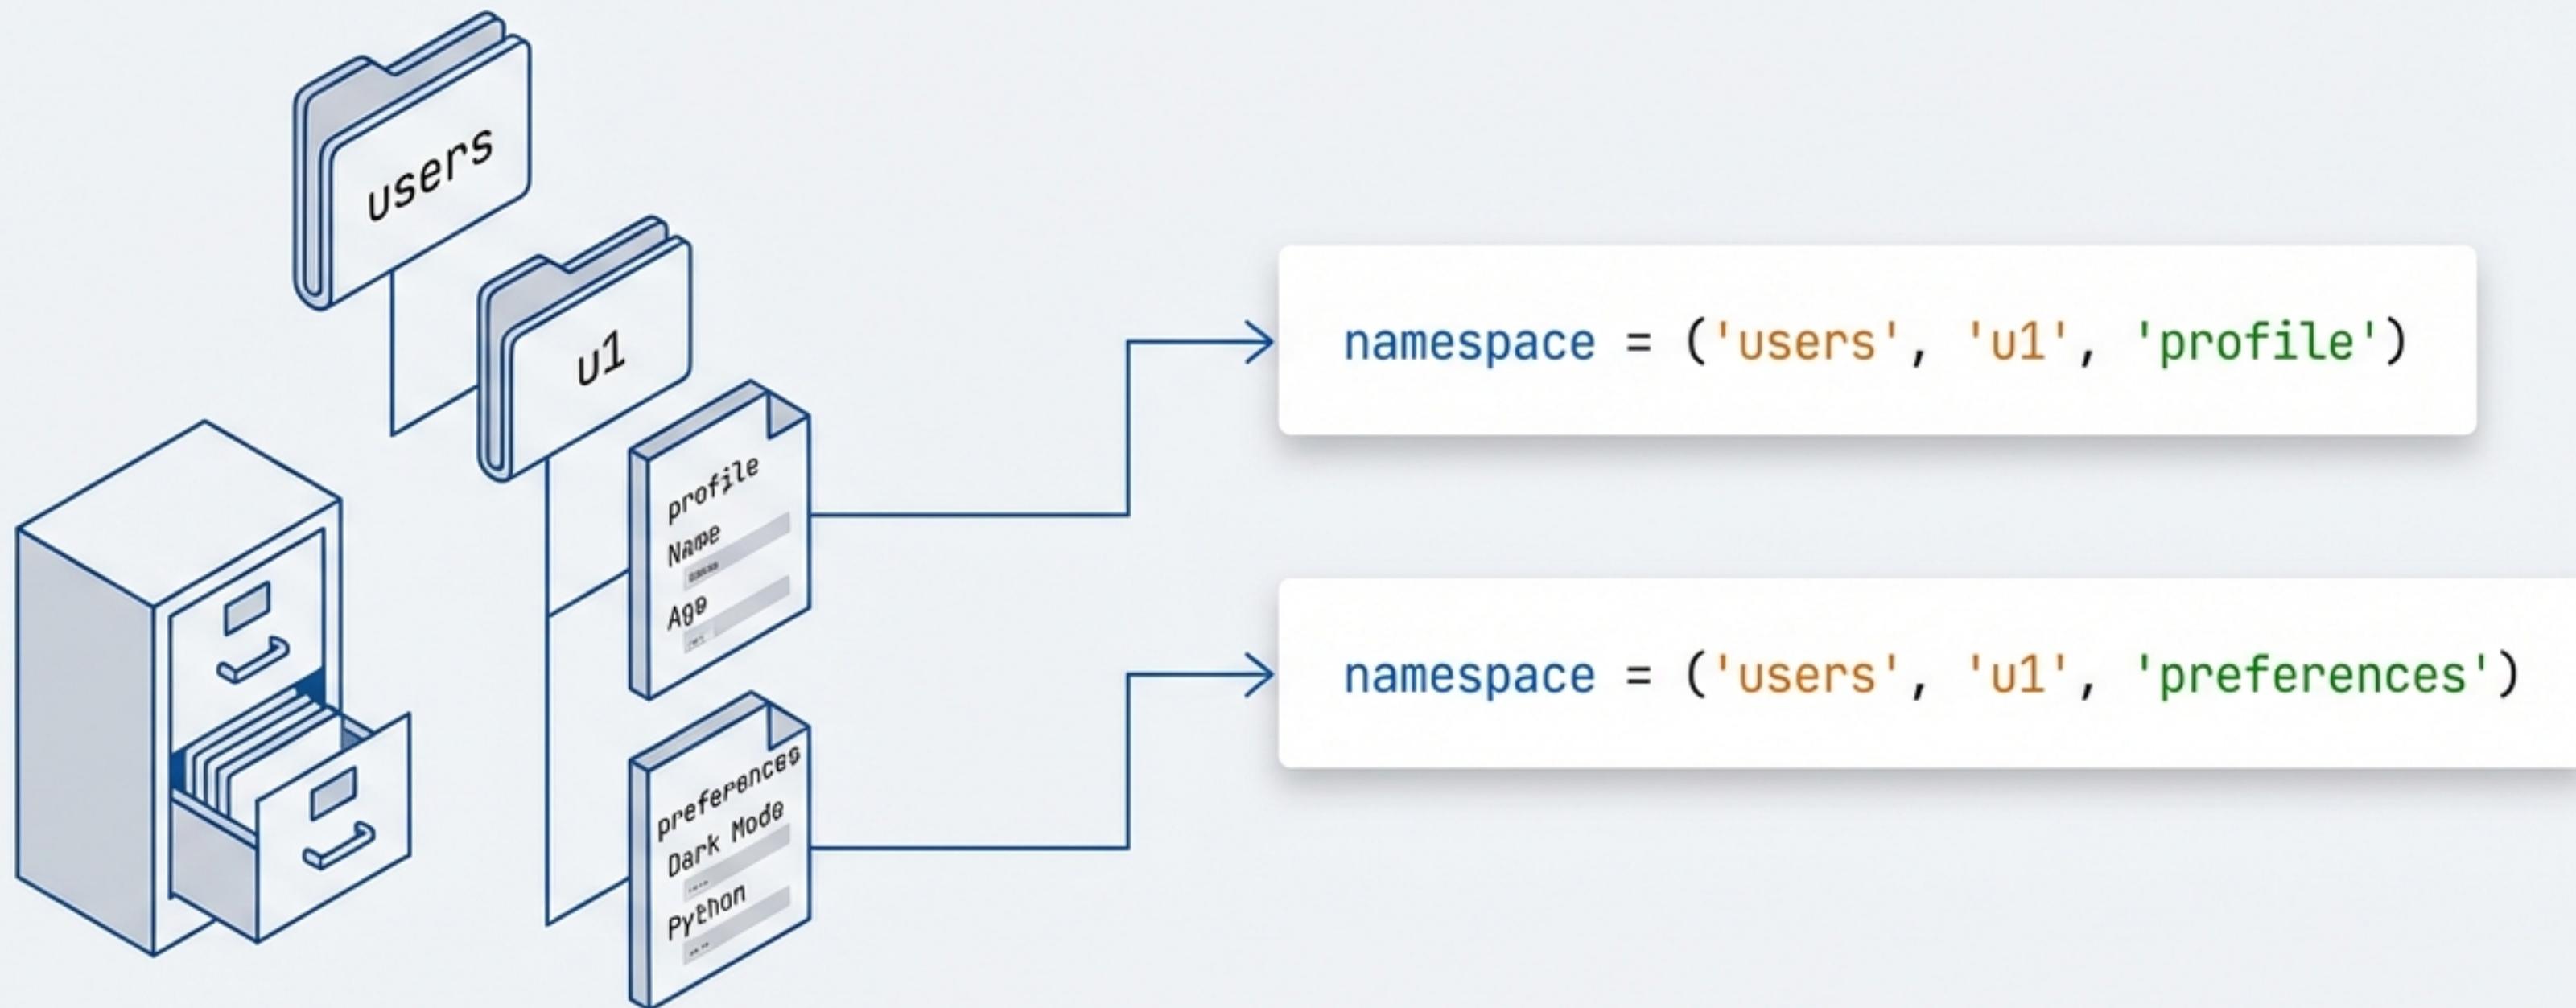
# The BaseStore Abstraction



```
from langgraph.store.memory import InMemoryStore
store = InMemoryStore()
```

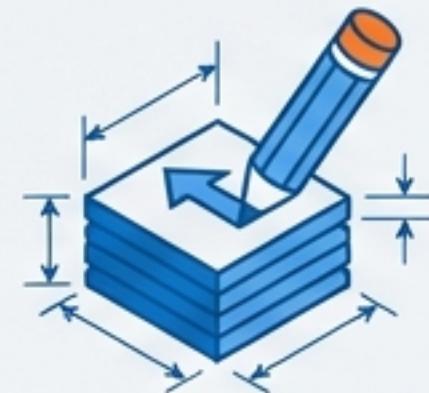
Stores data in RAM.  
Perfect for prototyping.

# Organizing Chaos with Namespaces



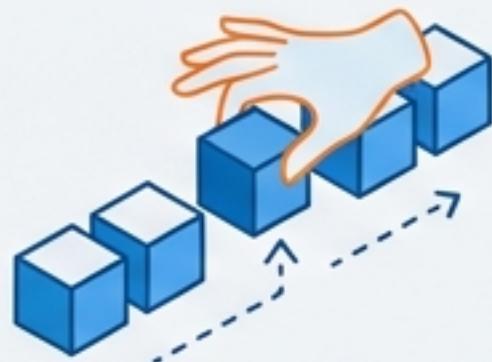
Namespaces are tuples that act as directory paths for memory isolation.

# The CRUD Operations



**PUT**  
(Write)

```
store.put(namespace, key='1',  
          value={'data': 'Likes Pizza'})
```



**GET**  
(Read Exact)

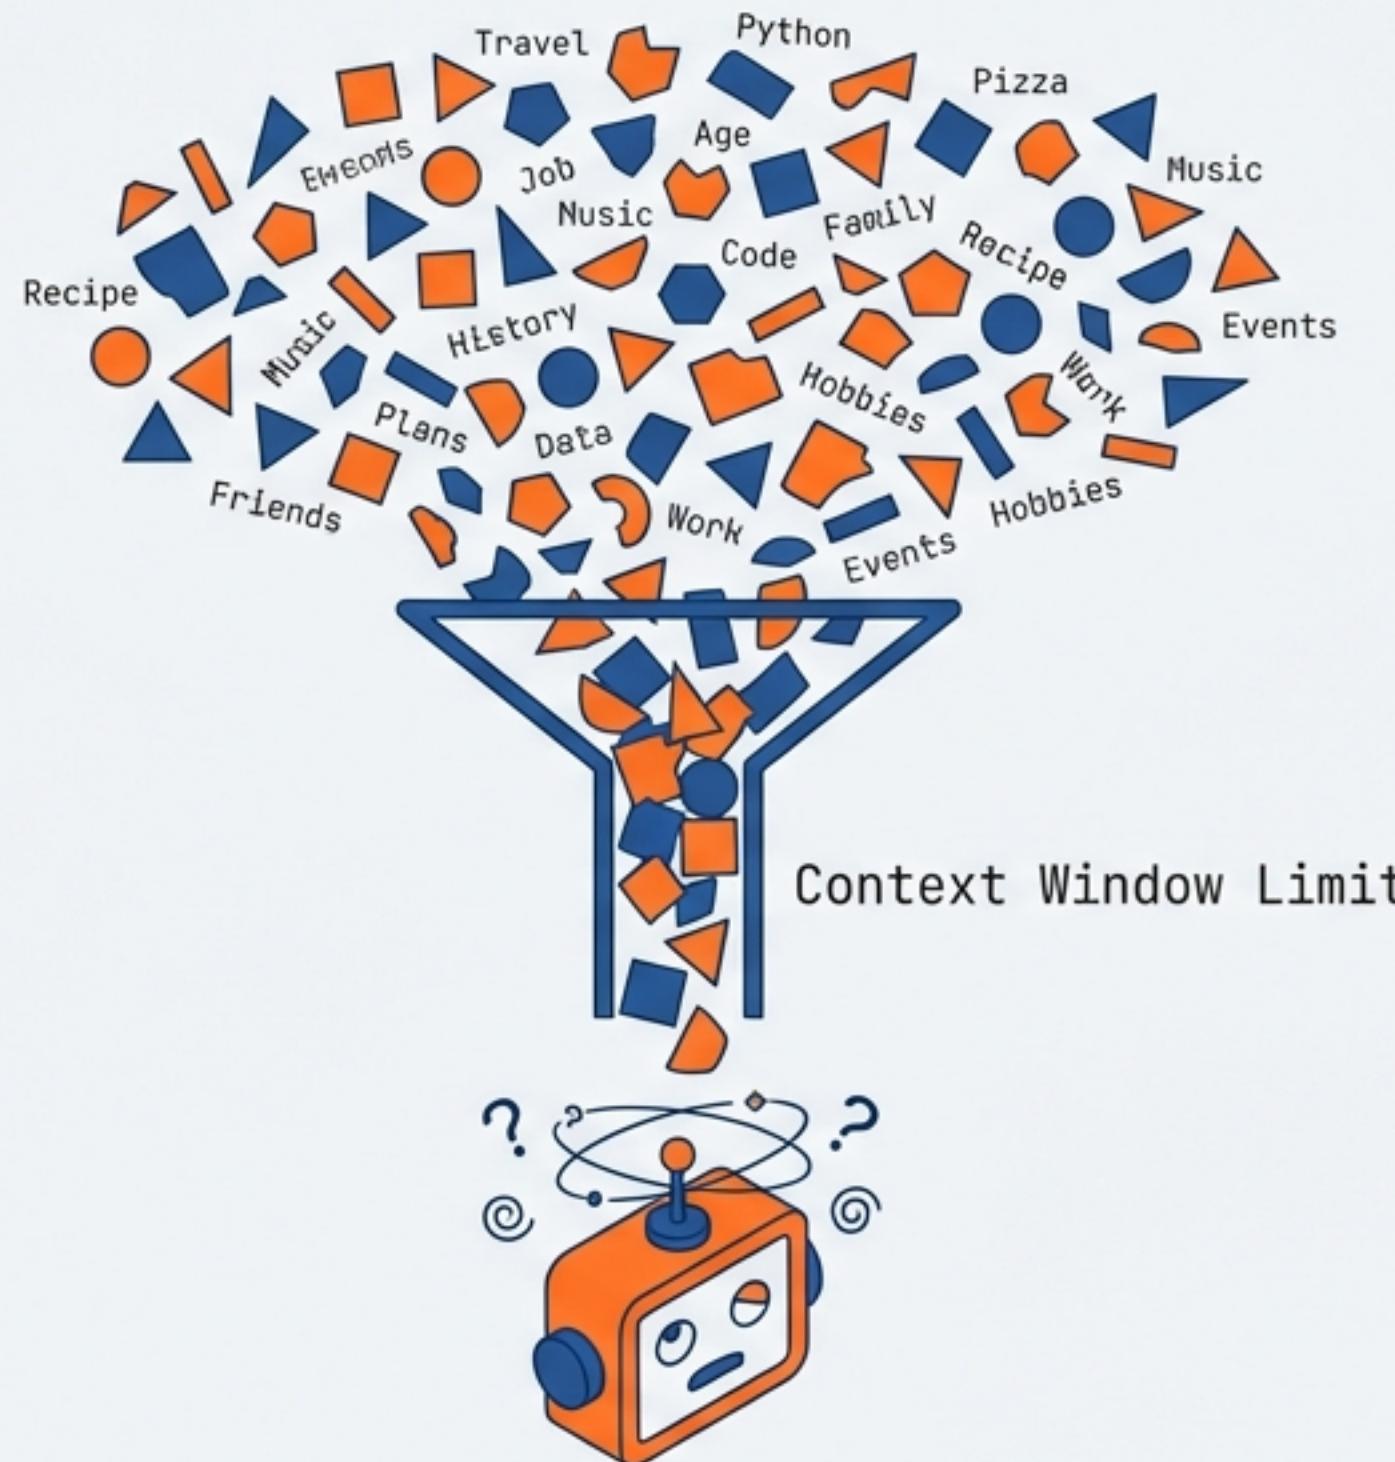
```
store.get(namespace, key='1')
```



**SEARCH**  
(Browse)

```
store.search(namespace)
```

# The Retrieval Bottleneck



Scenario:

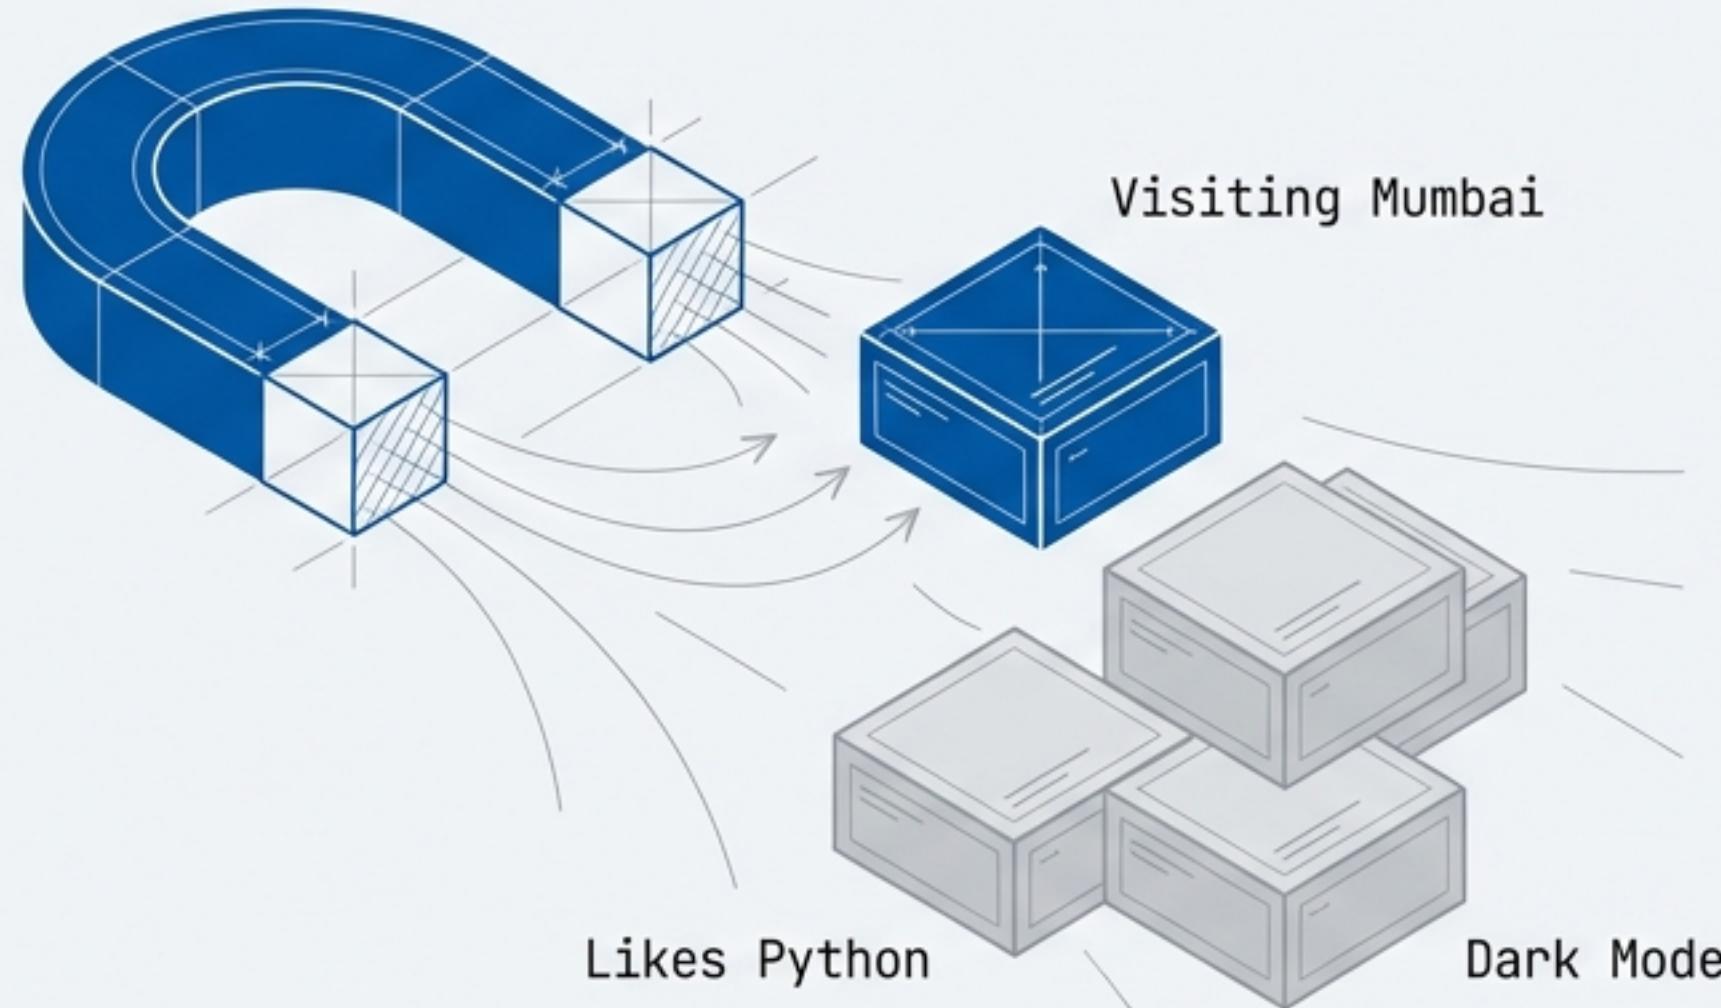
Query: "What is my travel plan?"

Insight:

Using `store.search()` dumps everything.  
We need a filter for meaning.

# Enabling Semantic Search

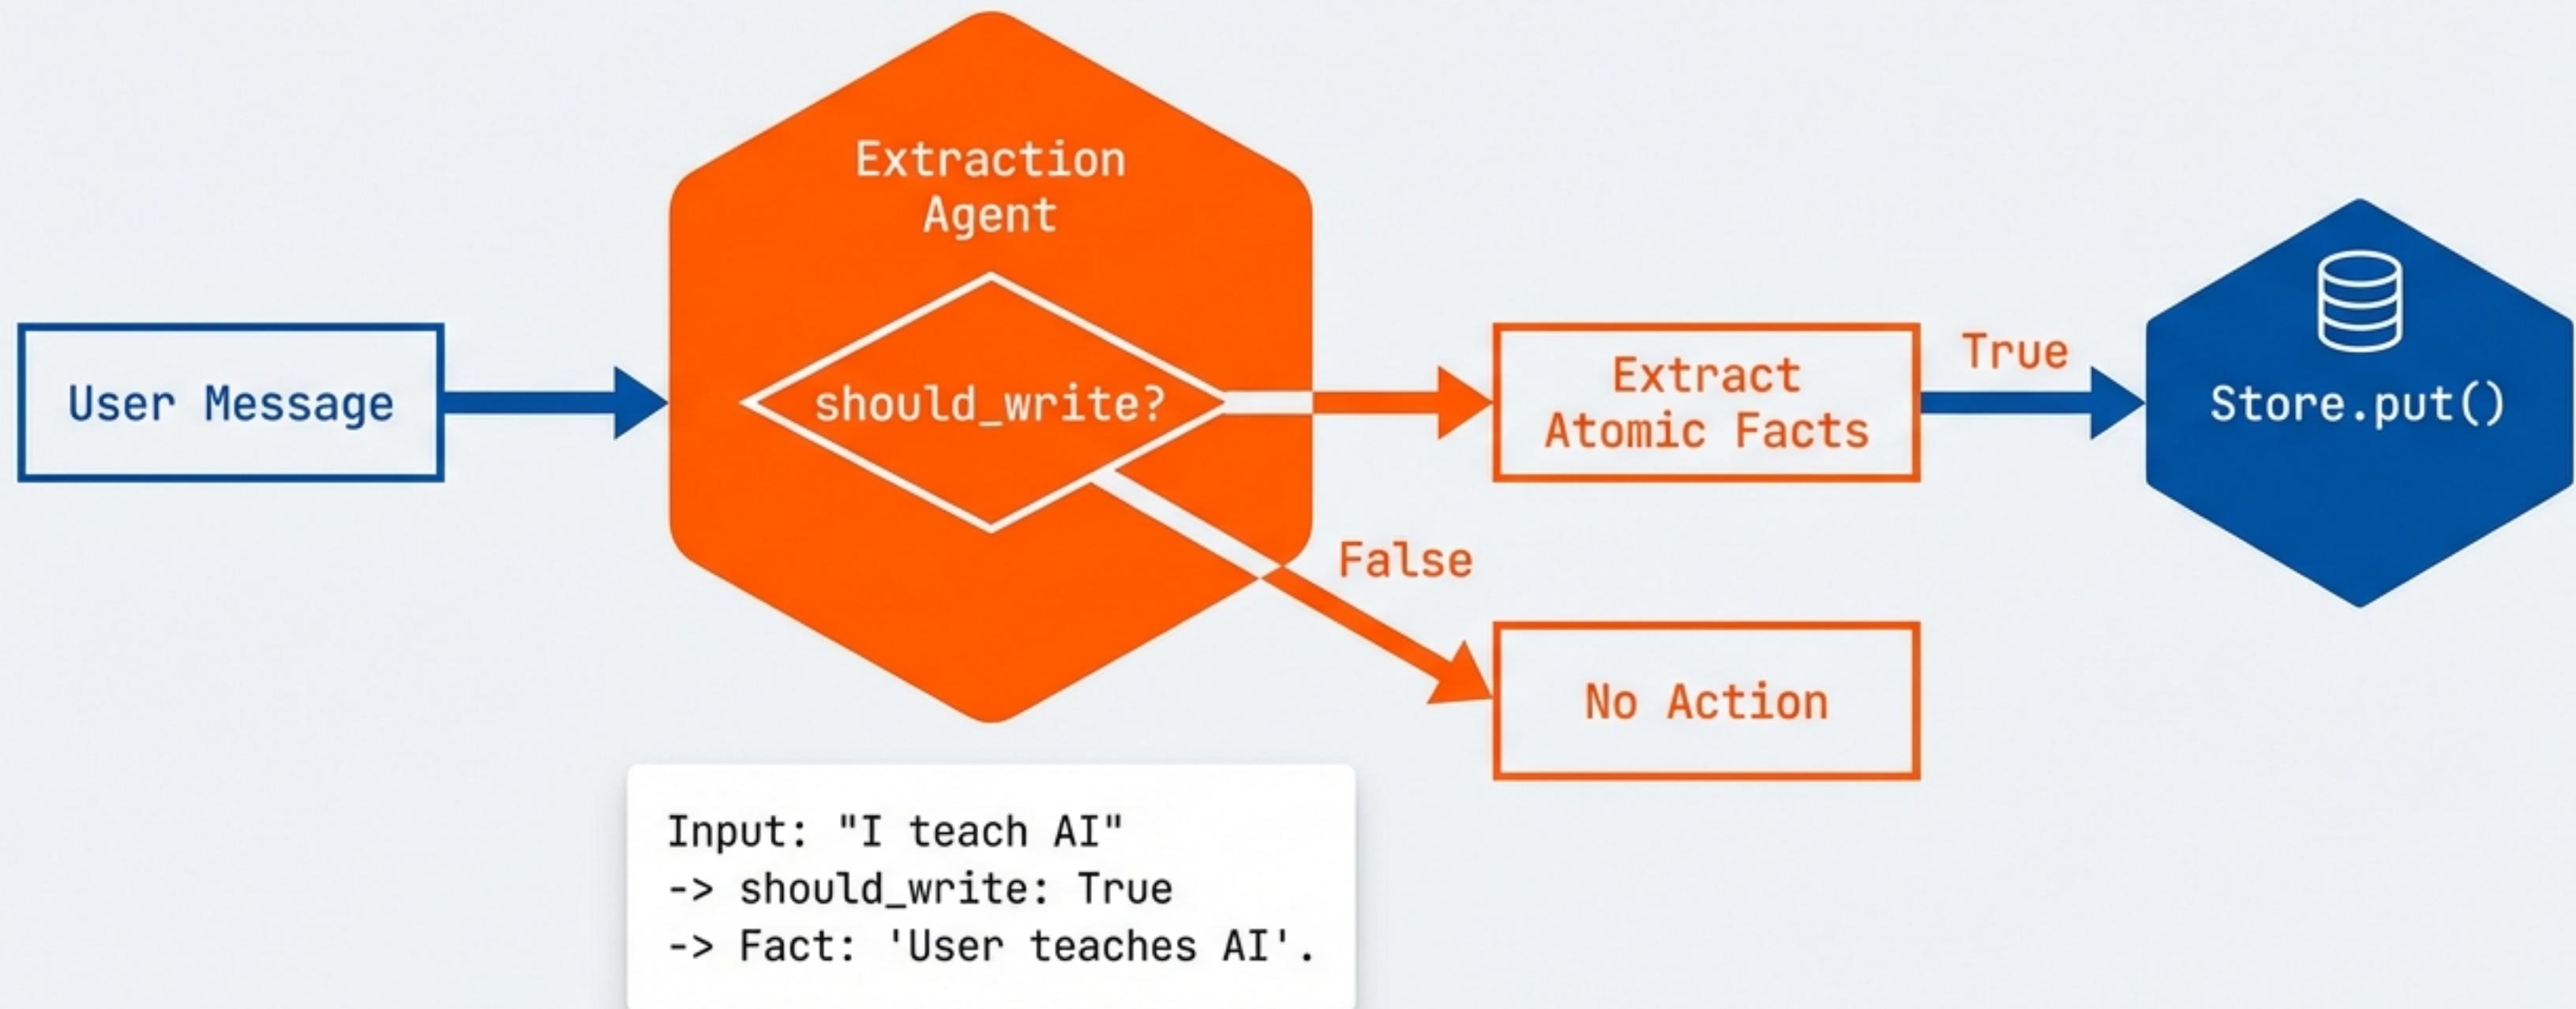
Query: "Travel plans"



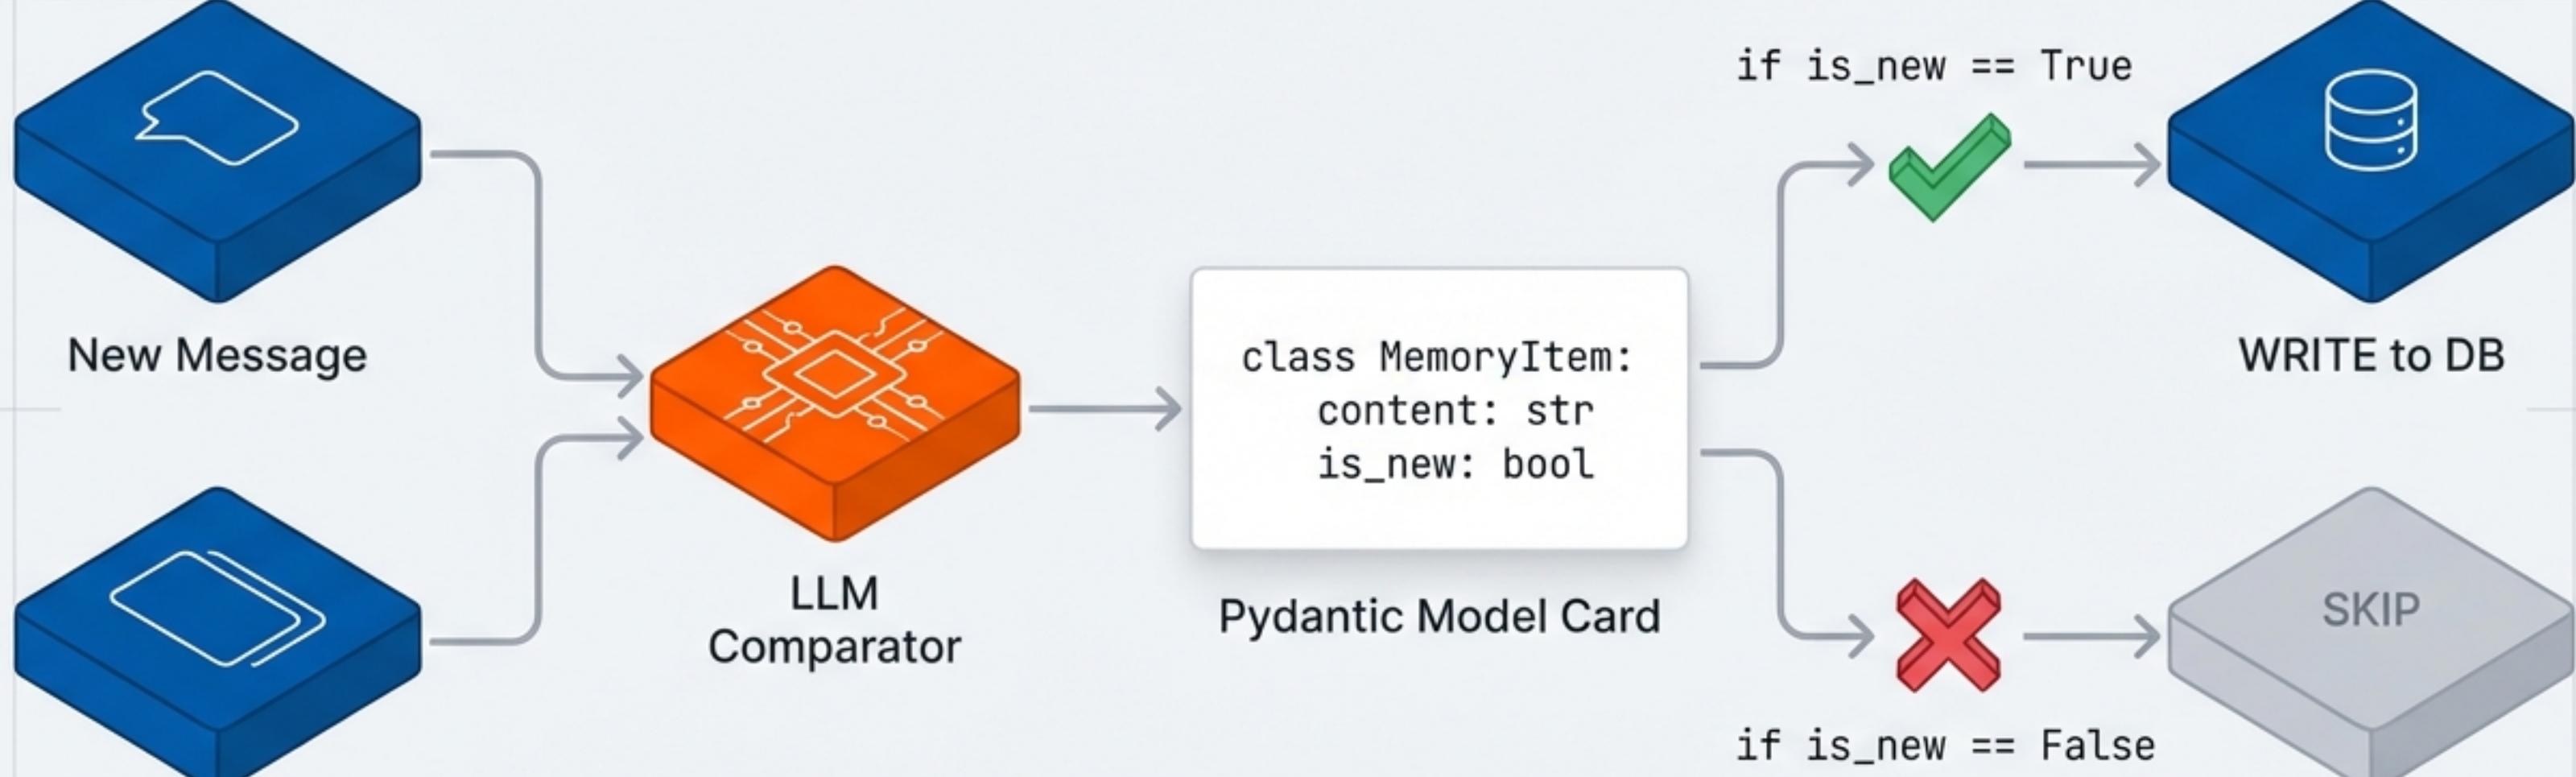
```
store.search(  
    namespace,  
    query='What is user currently learning?',  
    limit=1  
)
```

Vector embeddings allow retrieval by meaning, not just keywords.

# The Extraction Workflow

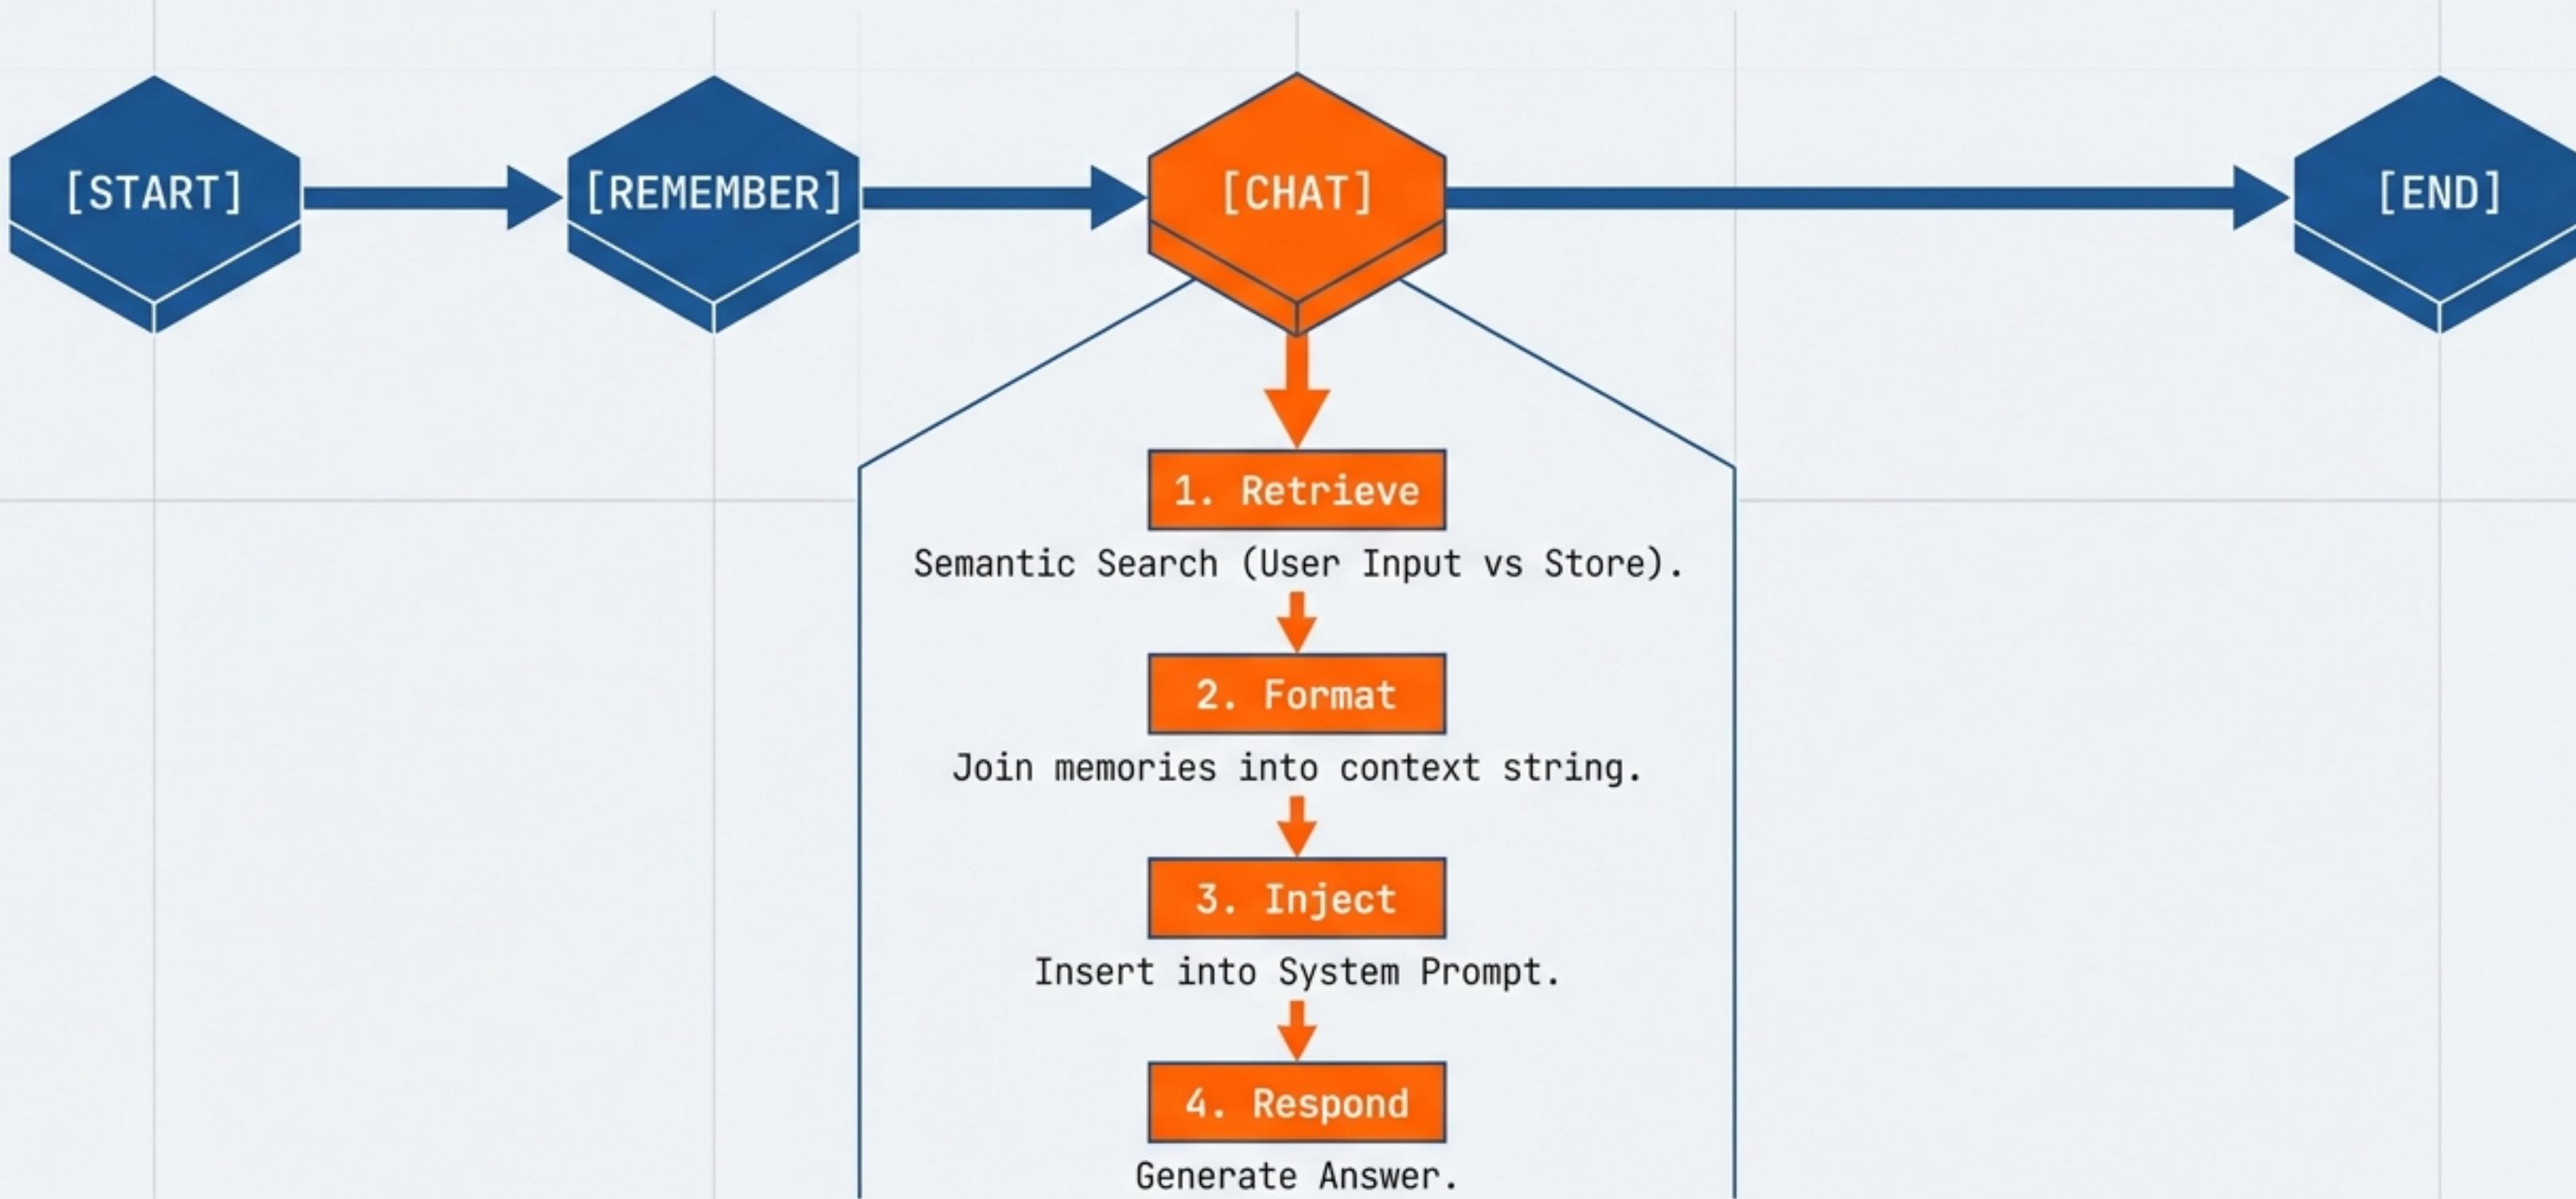


# Solving De-Duplication



Prevents storing redundant facts like 'My name is Nitish' multiple times.

# The Injection & Chat Workflow



# Engineering the System Prompt

You are a helpful assistant with **memory capabilities**.

Address user **by name**.

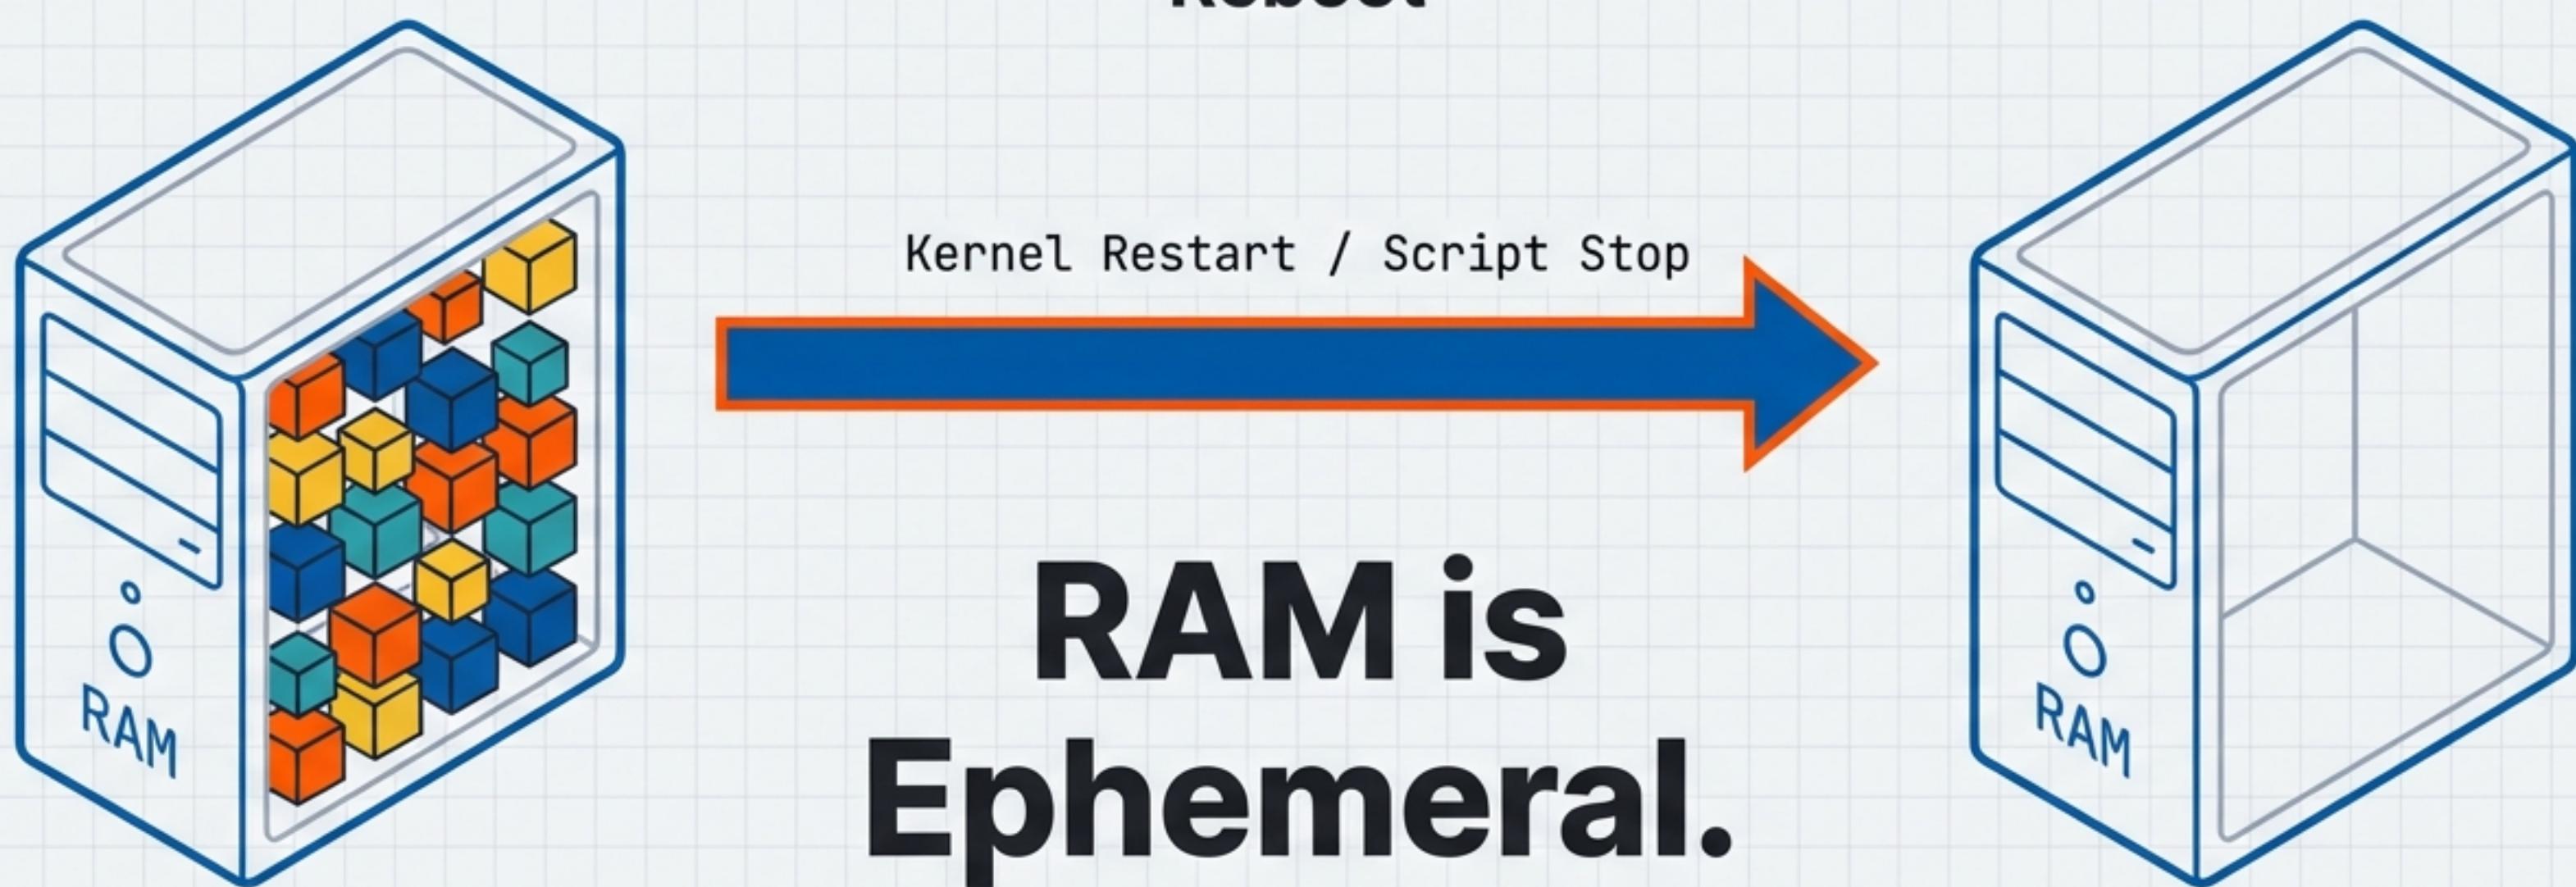
Reference known projects/preferences (e.g.,  
**'Since you use TypeScript...'**).

Adjust tone to be **friendly and natural**.

Suggest 3 relevant follow-up questions  
**based on profile**.

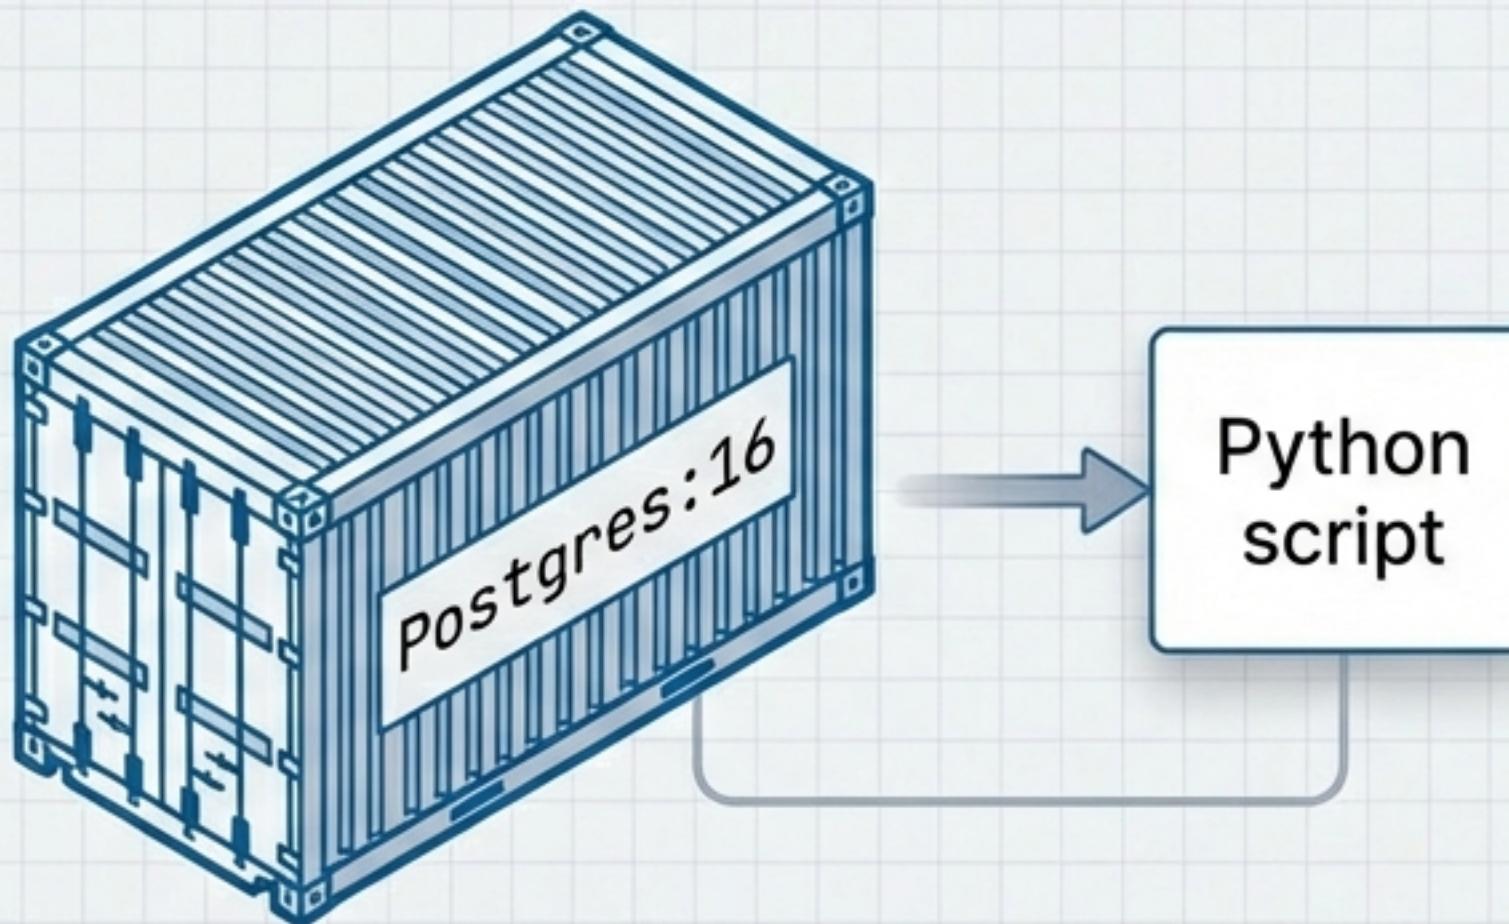
Forces the LLM to use  
the injected context.

# The Volatility of InMemoryStore



In-Memory prototyping is fast, but data vanishes on restart. Production requires persistence.

# Going Pro: The Postgres Store

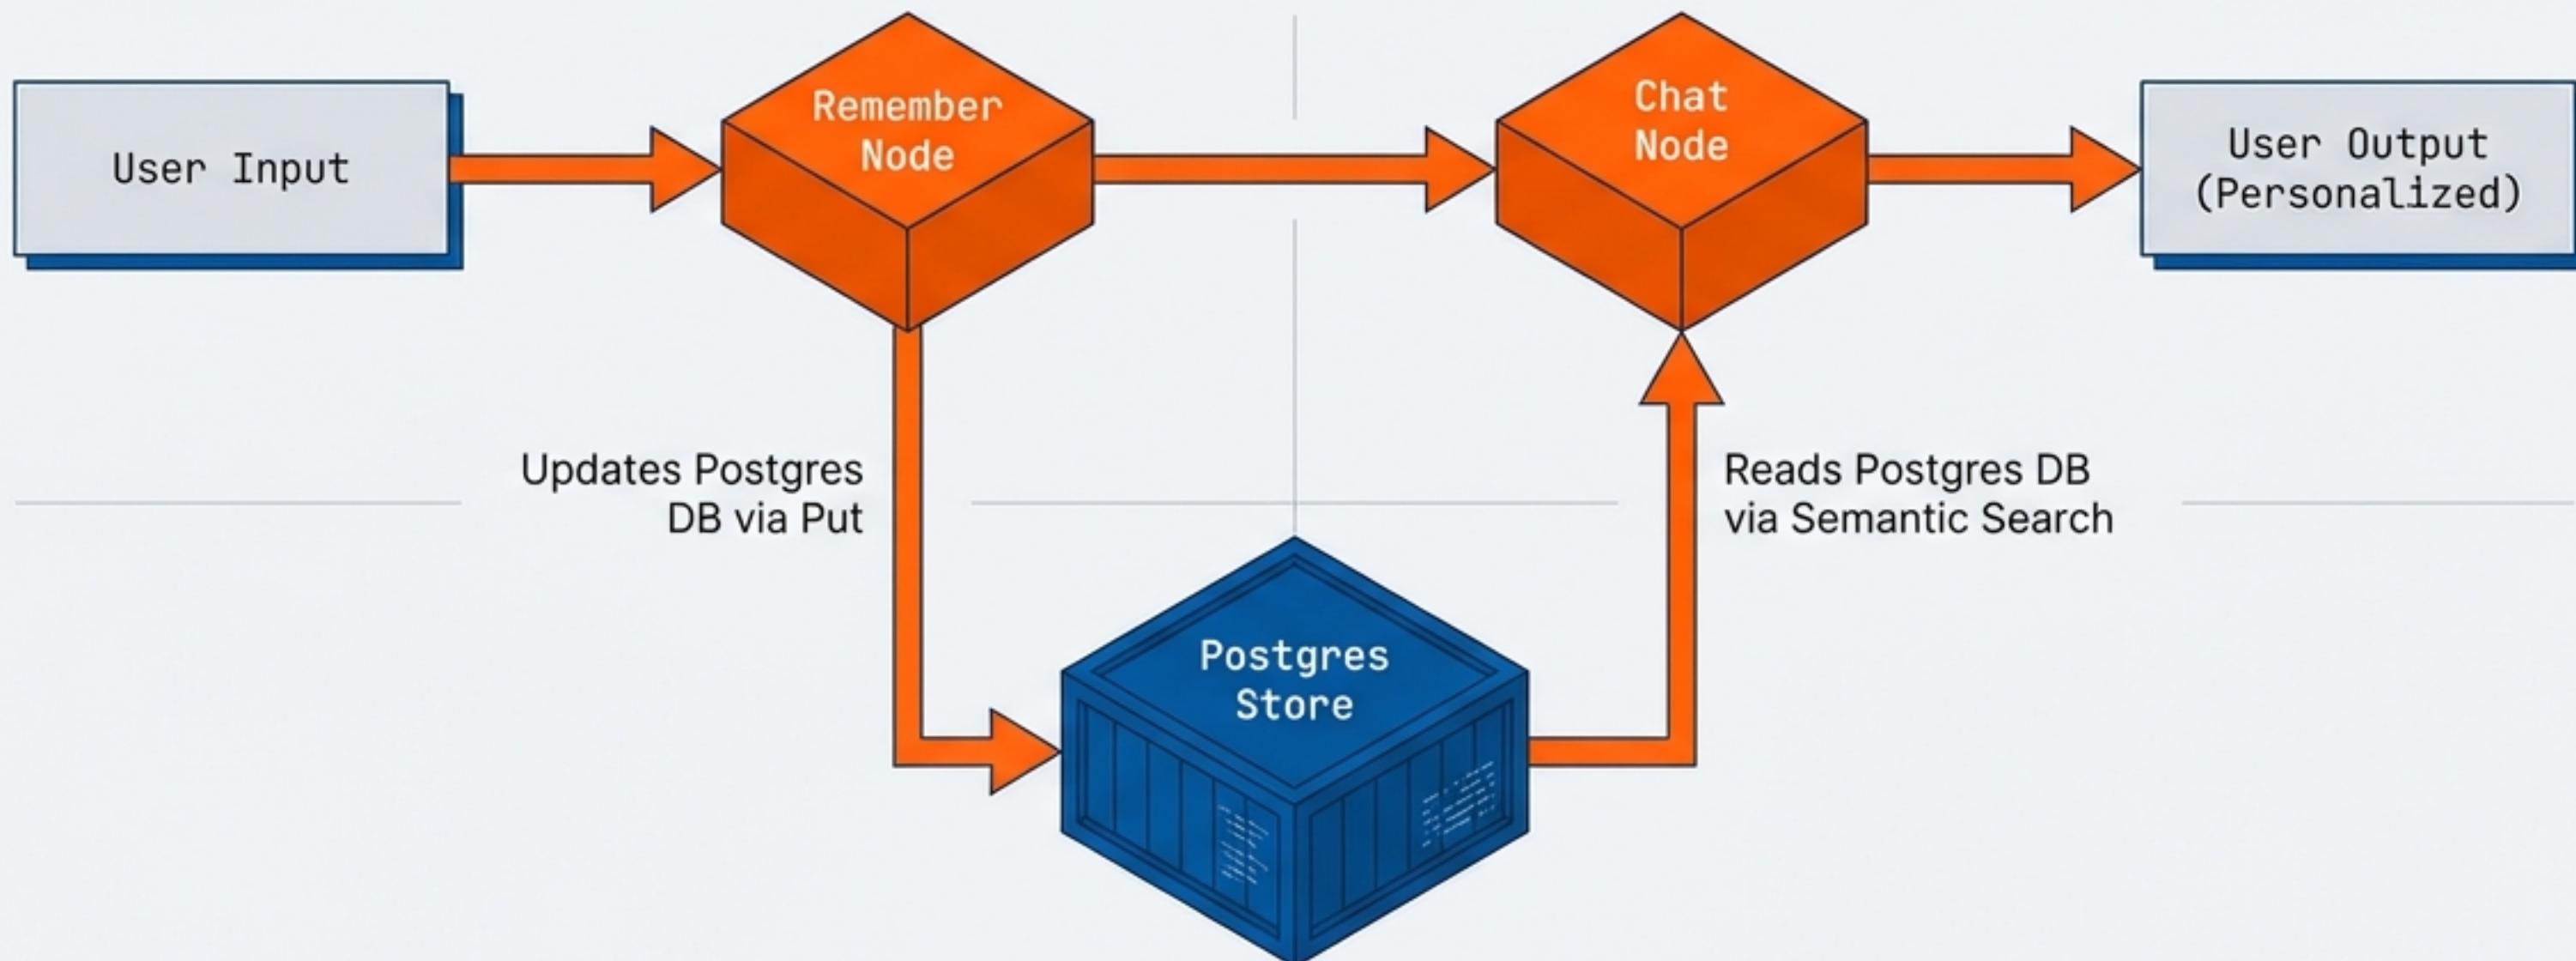


## CODE SNIPPET

```
- store = InMemoryStore()  
+ with  
    PostgresStore.from_conn_string  
        (DB_URL) as store:
```

The Graph logic remains unchanged. Only the backend is swapped.

# The Complete Persistent Agent



## CHAT EXAMPLE

User: Explain GenAI.

Bot: Sure Nitish, since you like Python, here is a code example...