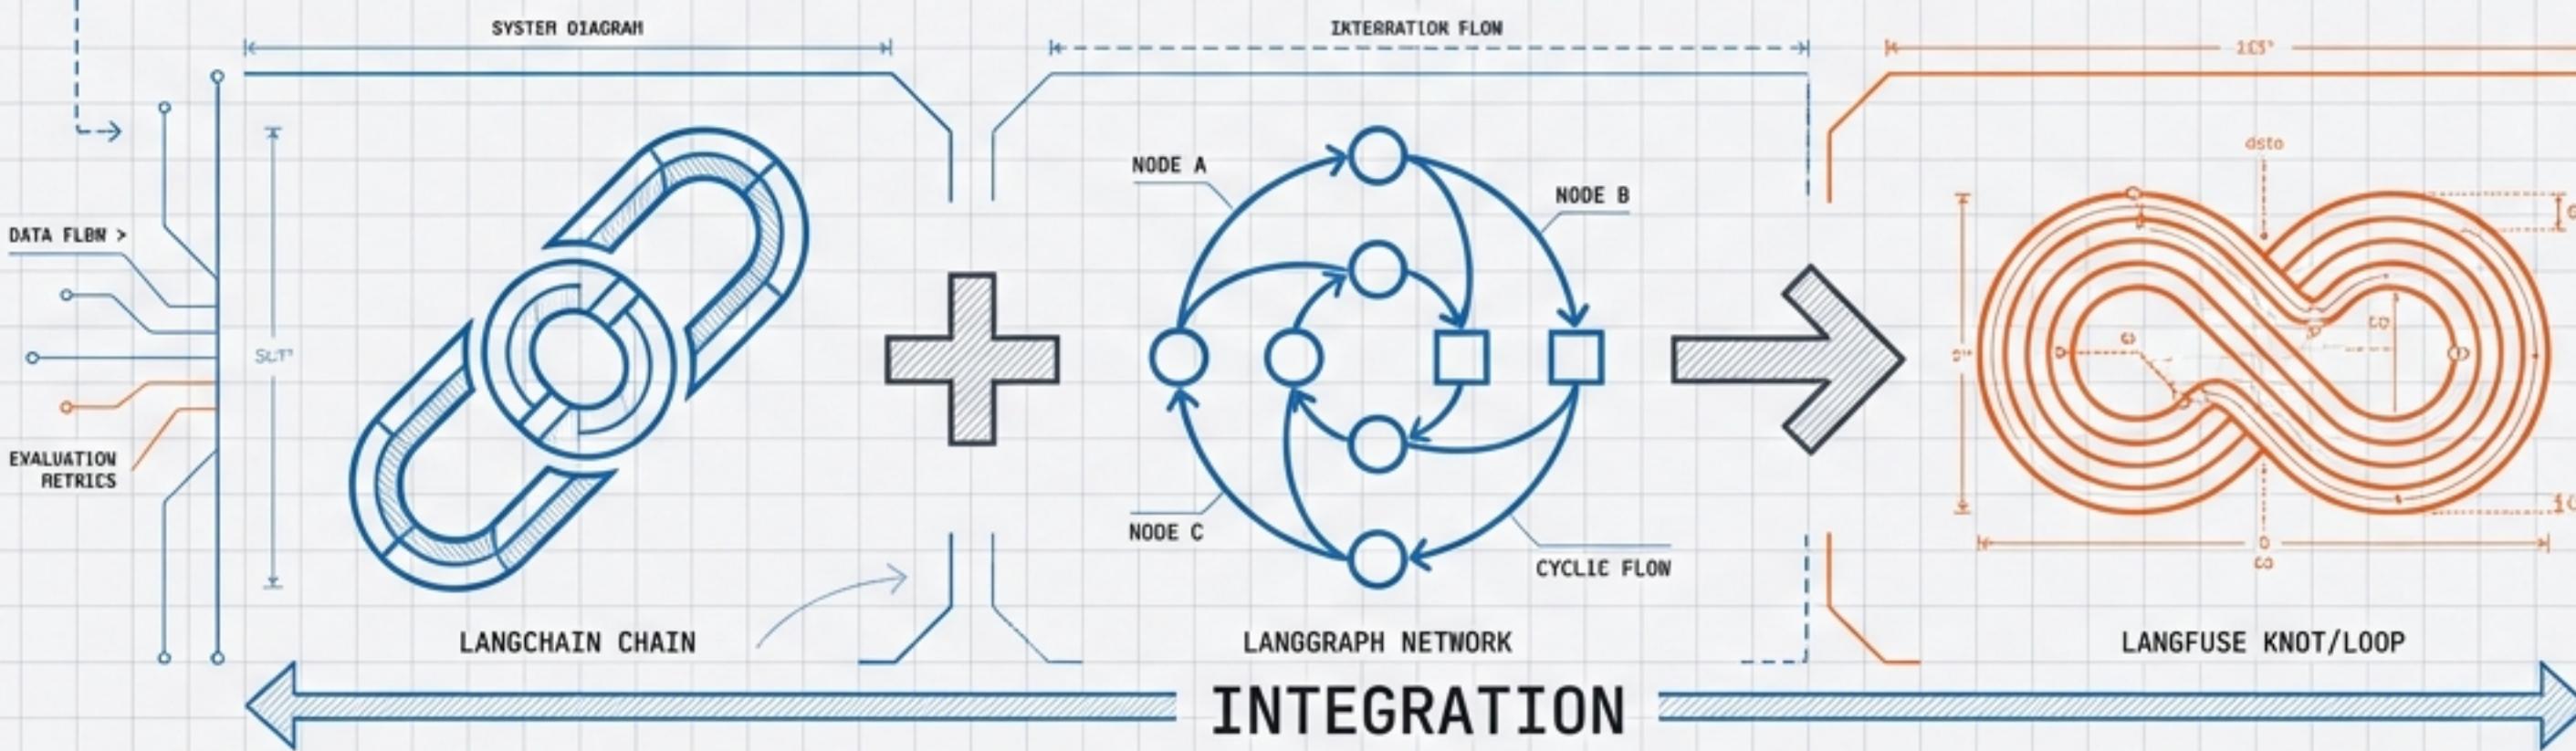
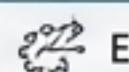


The Complete Guide to Langfuse for LangChain & LangGraph

Tracing, Evaluation, and Management for Python Applications



```
def trace_chain(data):
...
    langfuse.trace(
        langfuse.trace, ddes, ...)
)
```



ENGINEERING LIFECYCLE SERIES • VOL. 1

PROJECT: AI-OPS

DATE: OCT 26, 2024

ENGINEER: SUISSE NECTEN

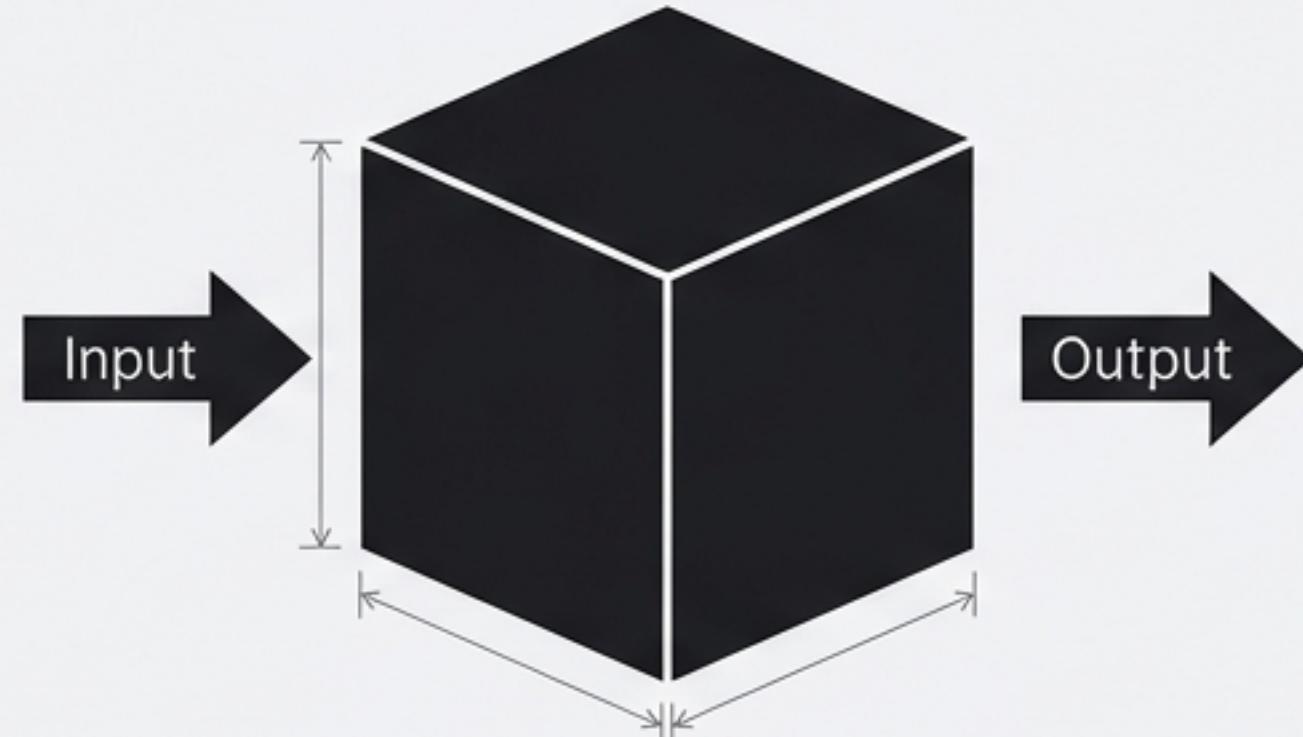
NotebookLM

The Observability Imperative

Moving from “it works on my machine”
to production reliability.

BEFORE

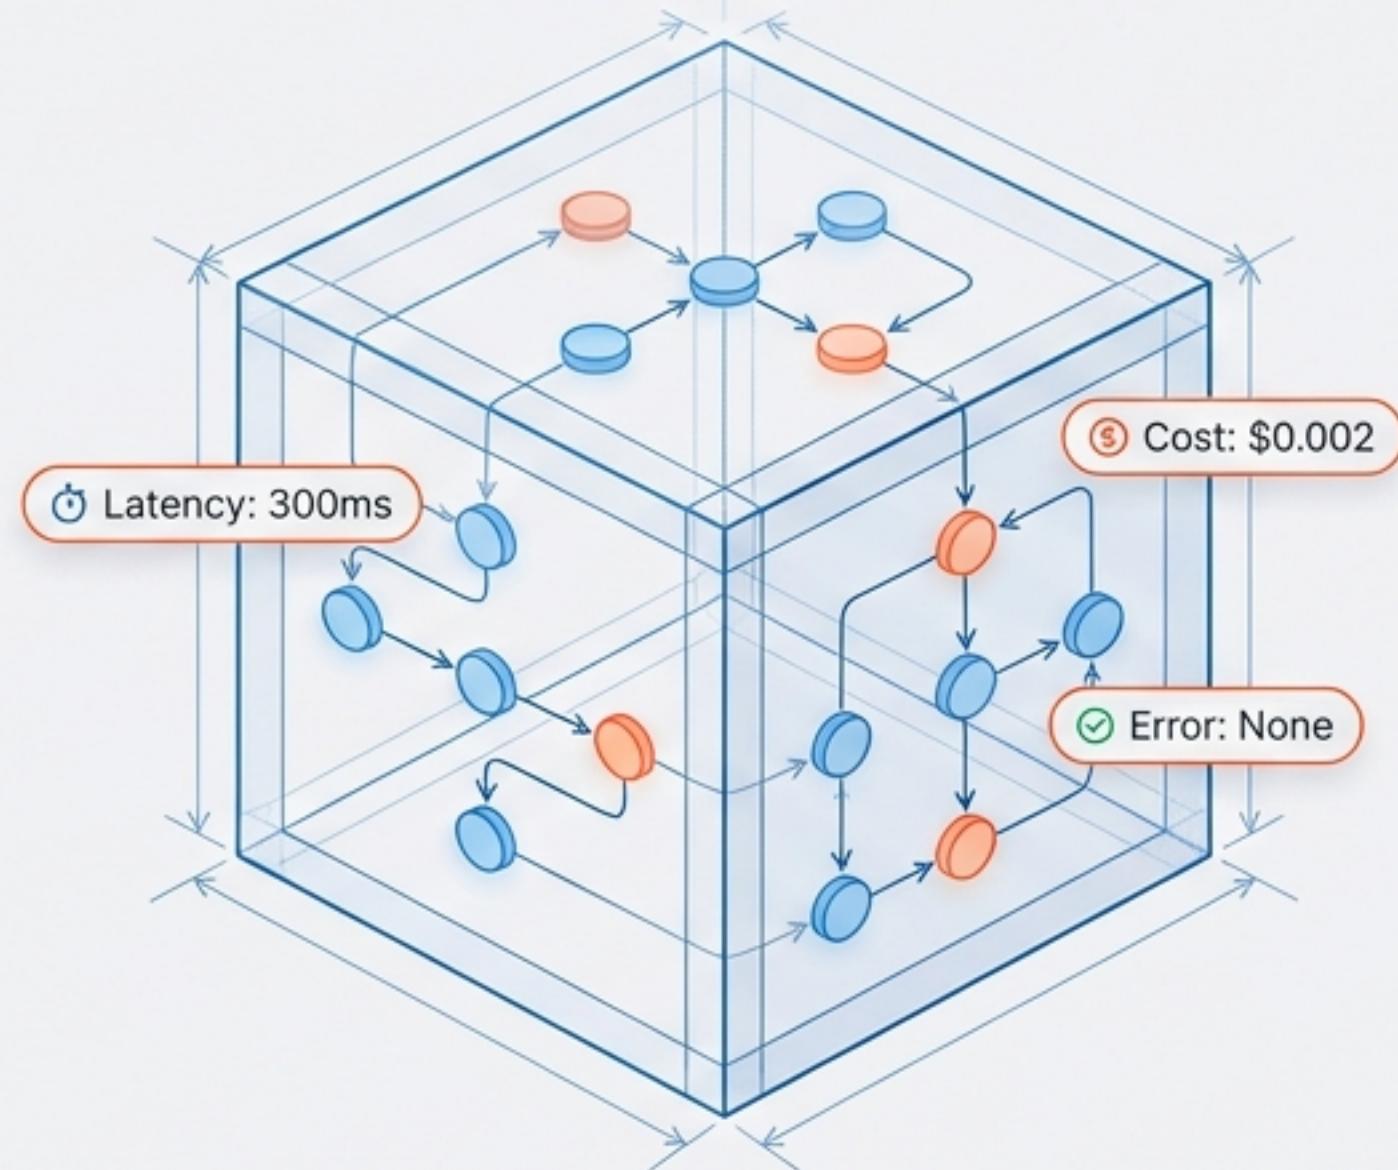
The Black Box



- ⌚ - Unclear failure points
- 💰 - Unknown costs
- ✳️ - Non-deterministic errors

AFTER

Glass Box



- ⌚ 1. DEBUG: Trace complex logic.
- ⌚ 2. MONITOR: Real-time health metrics.
- ☑ 3. IMPROVE: Dataset evaluation loops.

Initialization & Environment Setup

Step 01: Installation

```
%pip install langfuse langchain langgraph langchain_openai
```

Step 02: Environment Variables

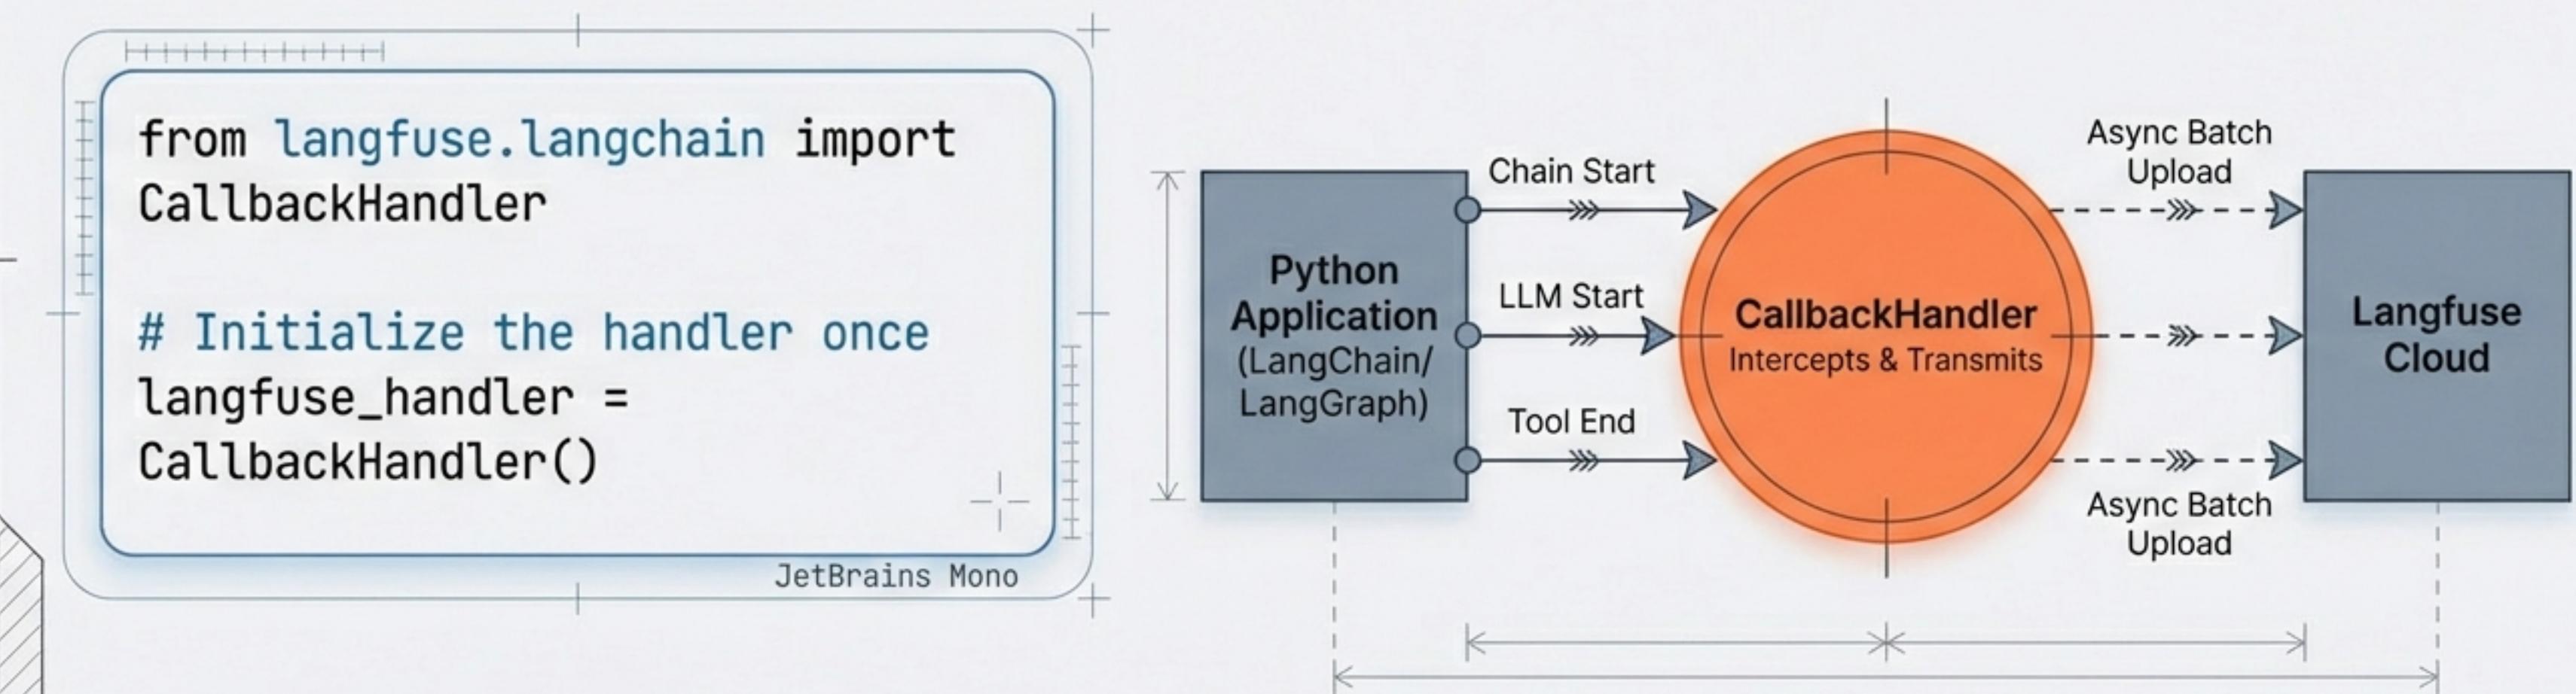
```
LANGFUSE_PUBLIC_KEY = pk-lf-...
LANGFUSE_SECRET_KEY = sk-lf-...
LANGFUSE_BASE_URL   = https://cloud.langfuse.com
OPENAI_API_KEY       = sk-proj-...
```

Step 03: Client Initialization

```
from langfuse import get_client
langfuse = get_client()
# Verify connection: langfuse.auth_check()
```

The Core Mechanism: CallbackHandler

The CallbackHandler is the bridge between your application logic and the Langfuse observability cloud. It intercepts events—Start, End, Error—and transmits them asynchronously.

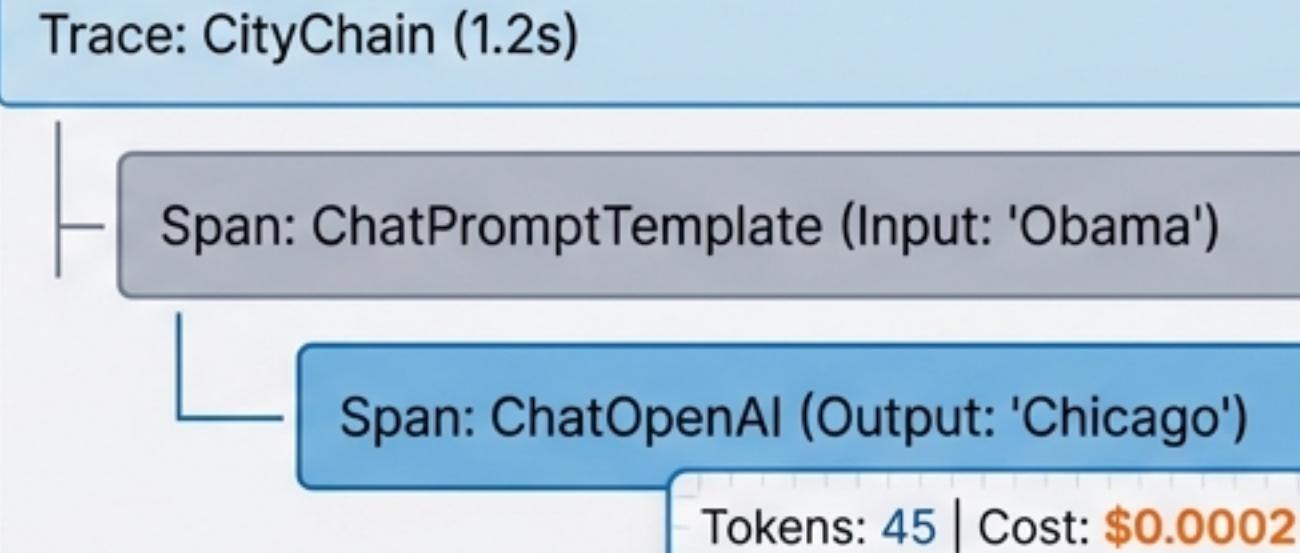


Tracing Linear Chains (LCEL)

THE INPUT (Python)

```
chain = prompt | model | StrOutputParser()  
  
# Pass the handler in the config  
chain.invoke(  
    {'person': 'Obama'},  
    config={'callbacks':[langfuse_handler]})
```

THE OUTPUT (Waterfall Trace)

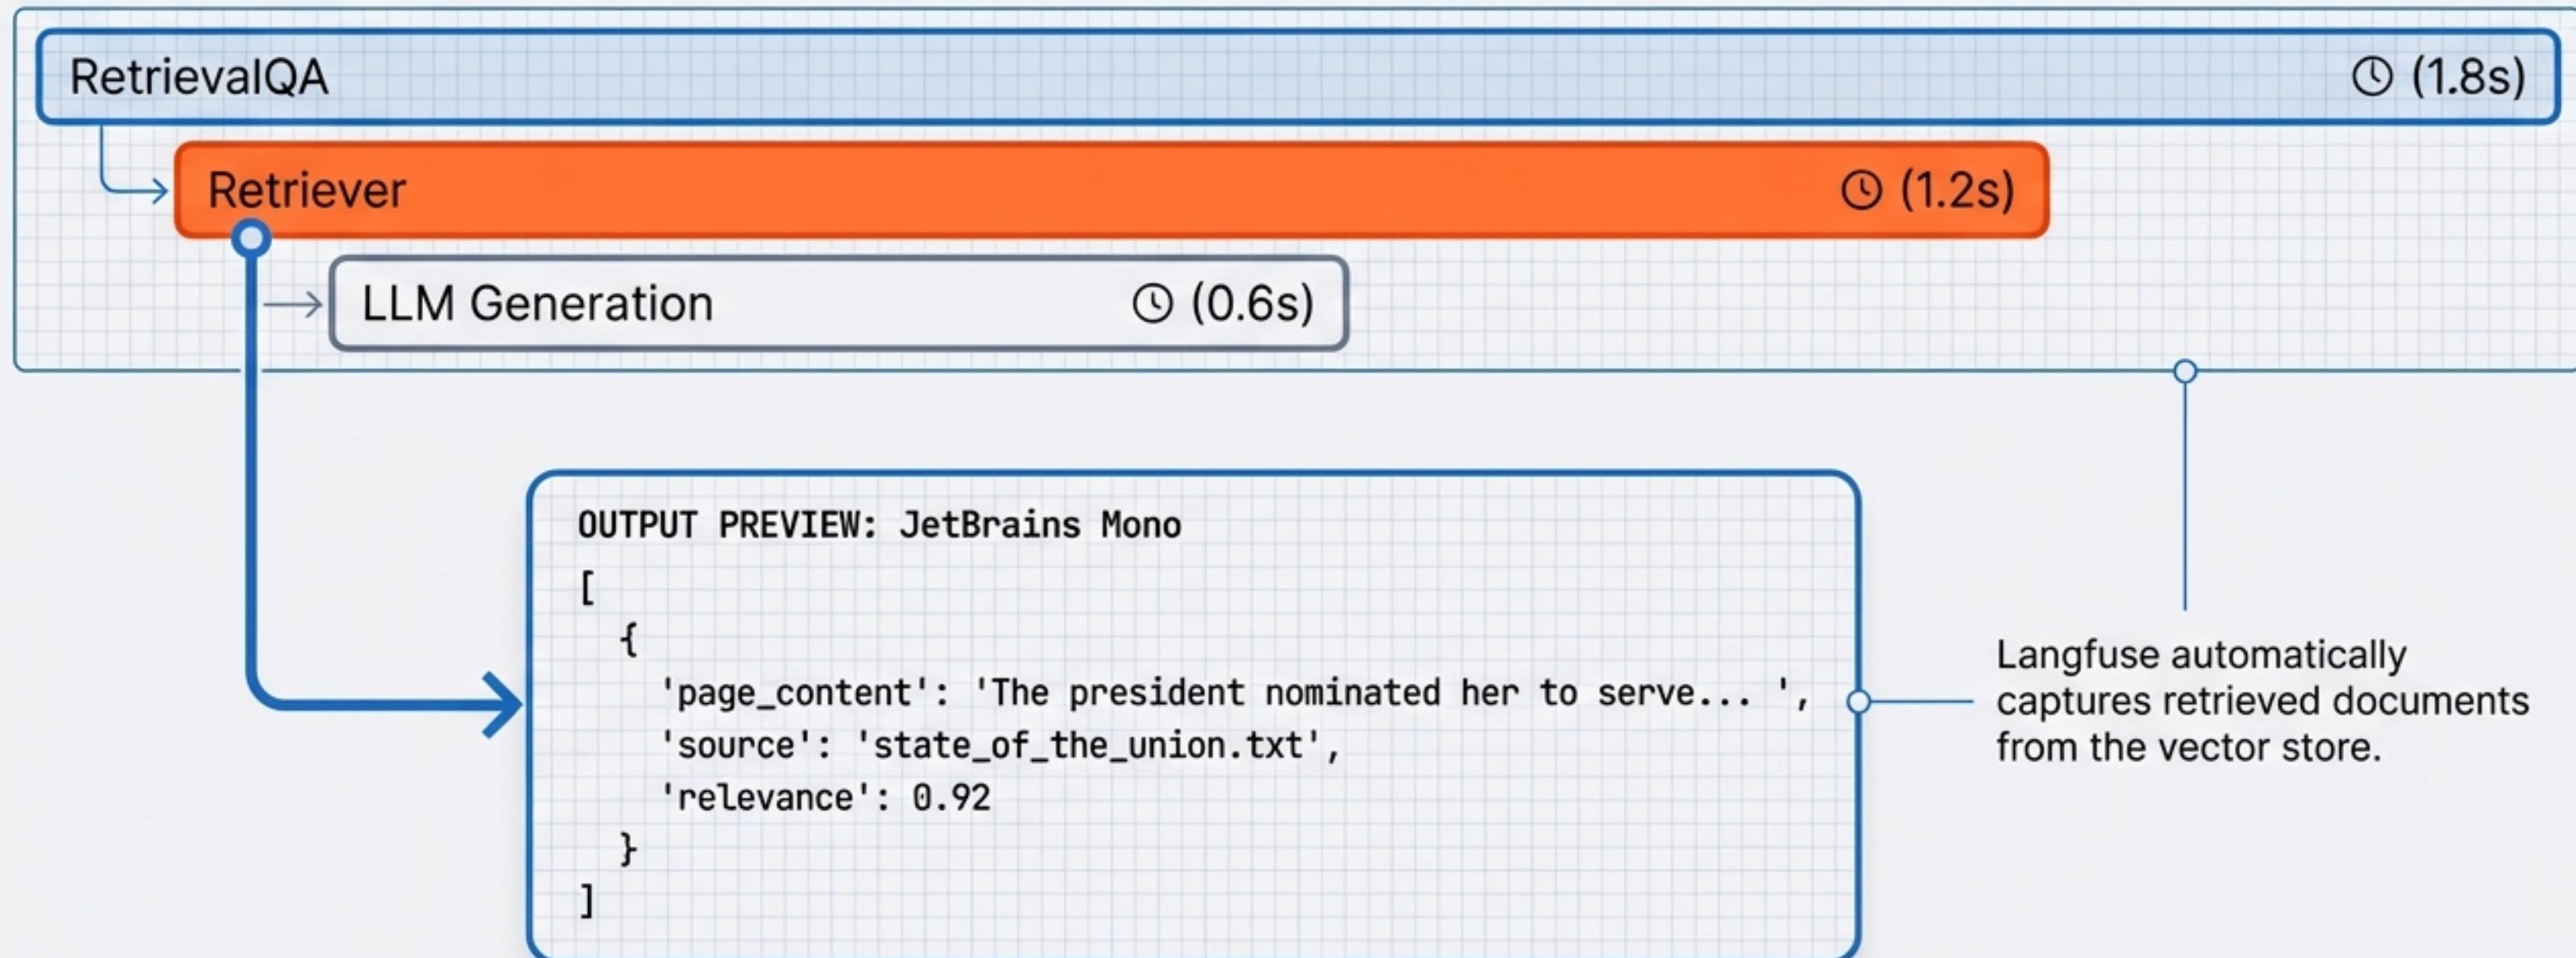


Debugging Retrieval Pipelines (RAG)

Inspecting retrieved context to identify hallucinations.

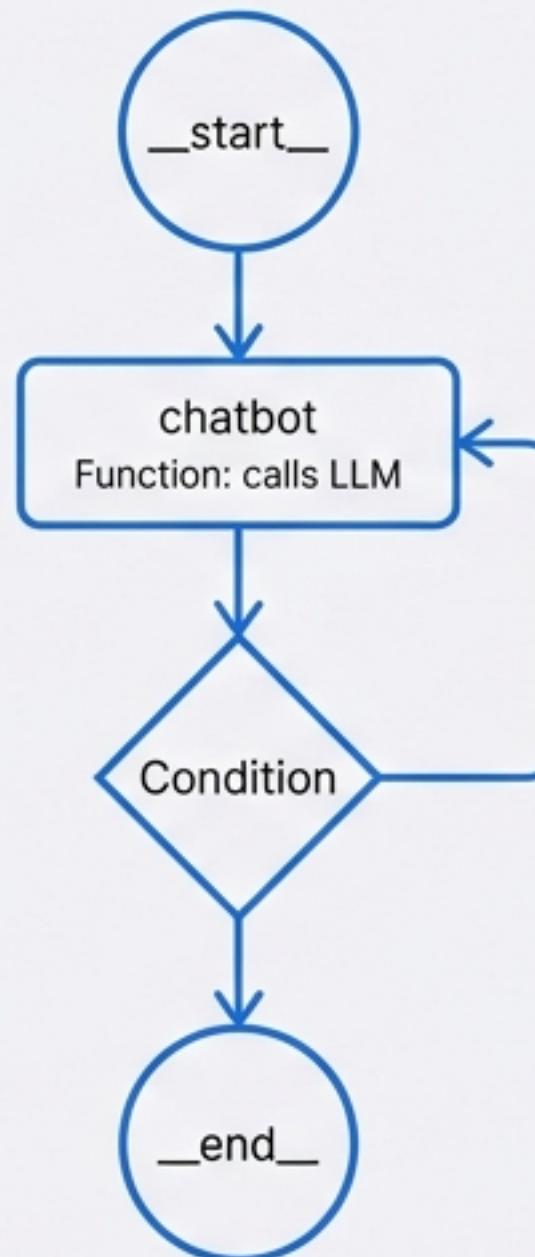
```
chain = RetrievalQA.from_chain_type(..., retriever=docsearch.as_retriever())
chain.invoke(query, config={'callbacks':[langfuse_handler]})
```

Trace Visualization



Stepping Up to LangGraph: Stateful Agents

Cyclic StateGraph



Implementation details

```
# 1. Compile the graph  
graph = builder.compile()  
  
# 2. Stream with callbacks  
inputs = {'messages': [HumanMessage(content='Hi')]}  
graph.stream(  
    inputs,  
    config={'callbacks': [langfuse_handler]})
```

The handler automatically traces the graph structure, preserving the parent/child relationship of steps.

Orchestrating Multi-Agent Systems

Supervisor Agent (LLM)

Decides next step: [Researcher, CurrentTime, FINISH]

Researcher

Tool: WikipediaQueryRun

CurrentTime

Tool: datetime_tool

Trace Tree

```
▼ 🤖 Supervisor (Trace Root)
  └─ Router (LLM Decision: 'Researcher')
    └─ 🔧 Researcher (Tool Execution)
      └─ ➤ wikipedia.search('Photosynthesis')
    └─ 🤖 Supervisor (Next Step: 'FINISH')
```

Langfuse visualizes the delegation of tasks across agent nodes.

Advanced Tracing: Linking & Custom Spans

Linking Distributed Executions



```
trace_id = Langfuse.create_trace_id()  
...  
trace_context={'trace_id': trace_id}
```

Manual Instrumentation (Custom Spans)

```
with  
langfuse .start_as_current_observation(  
    name='tv-sub-research-agent',  
    trace_context=...  
) as span:  
    # Execute custom logic here  
    result = sub_agent.invoke(...)  
    span.update_trace(output=result)
```

Wraps logic inside a custom span, useful for tools or sub-agents.

Online Evaluation (Production Metrics)

Monitoring the three vital signs of AI applications.

Costs (Token Usage)

Automatic Tracking



Latency (Bottlenecks)

Step-level breakdown

Retrieval (2.1s)

Generation (0.8s)

Explicit User Feedback

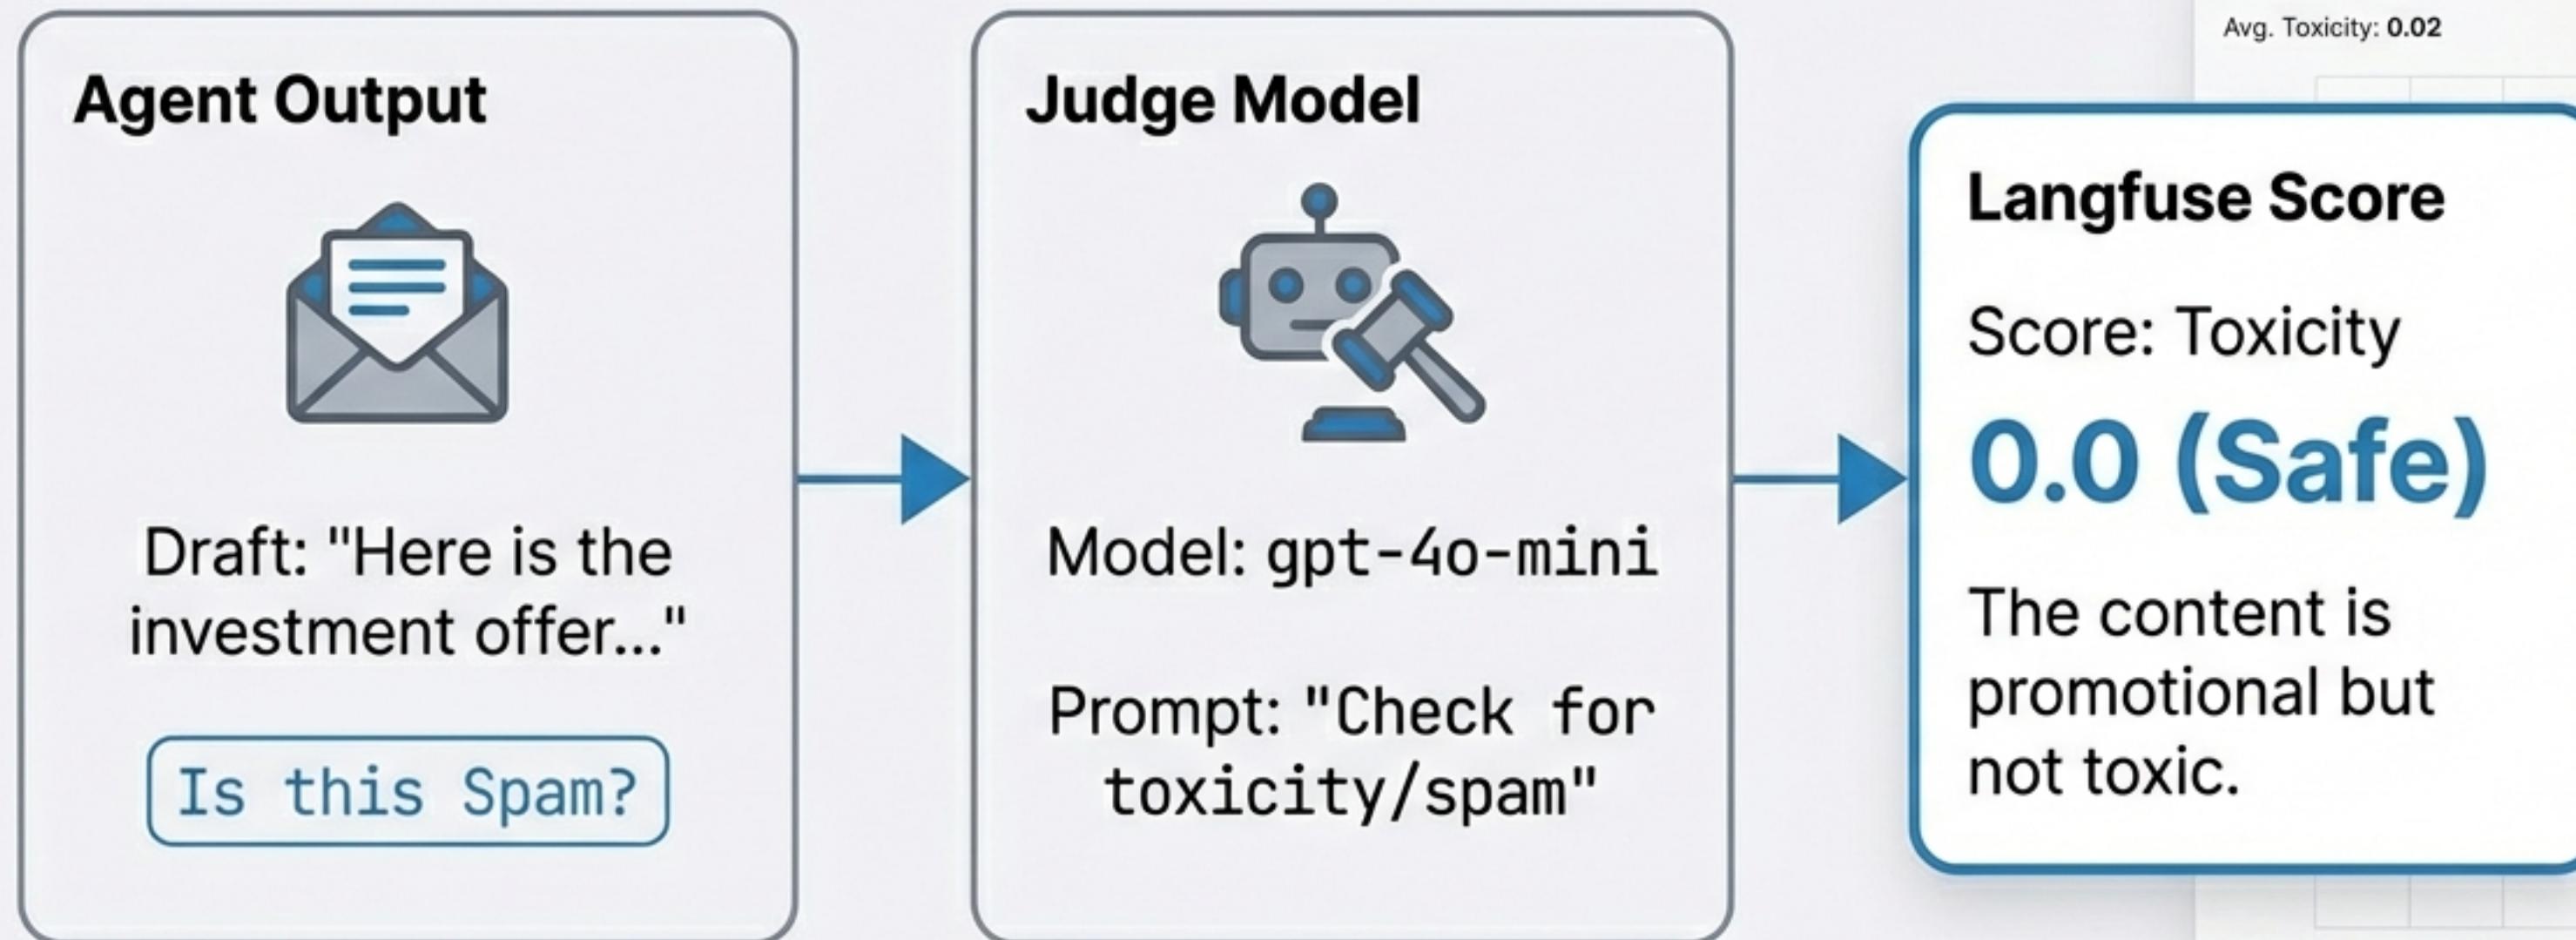


Correct answer

```
span.score_trace(  
    name='user-feedback',  
    value=1,  
    comment='Correct answer'  
)
```

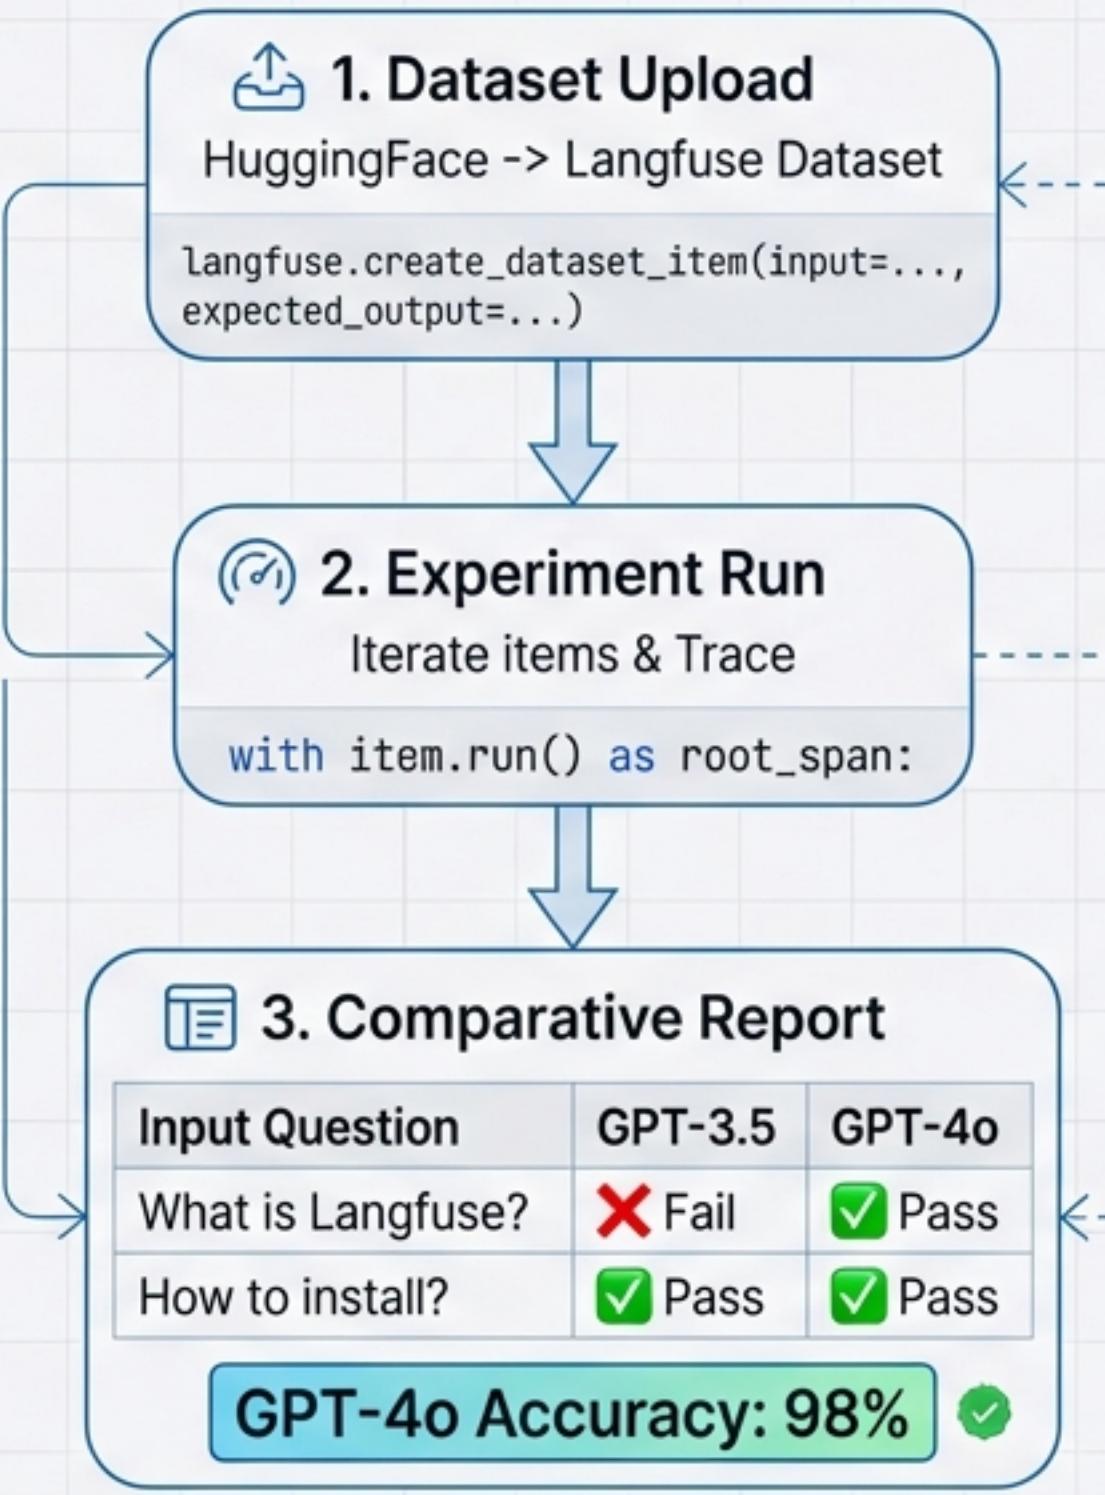
Automated QA: LLM-as-a-Judge

Scalable evaluation without human intervention.



Offline Evaluation (Datasets)

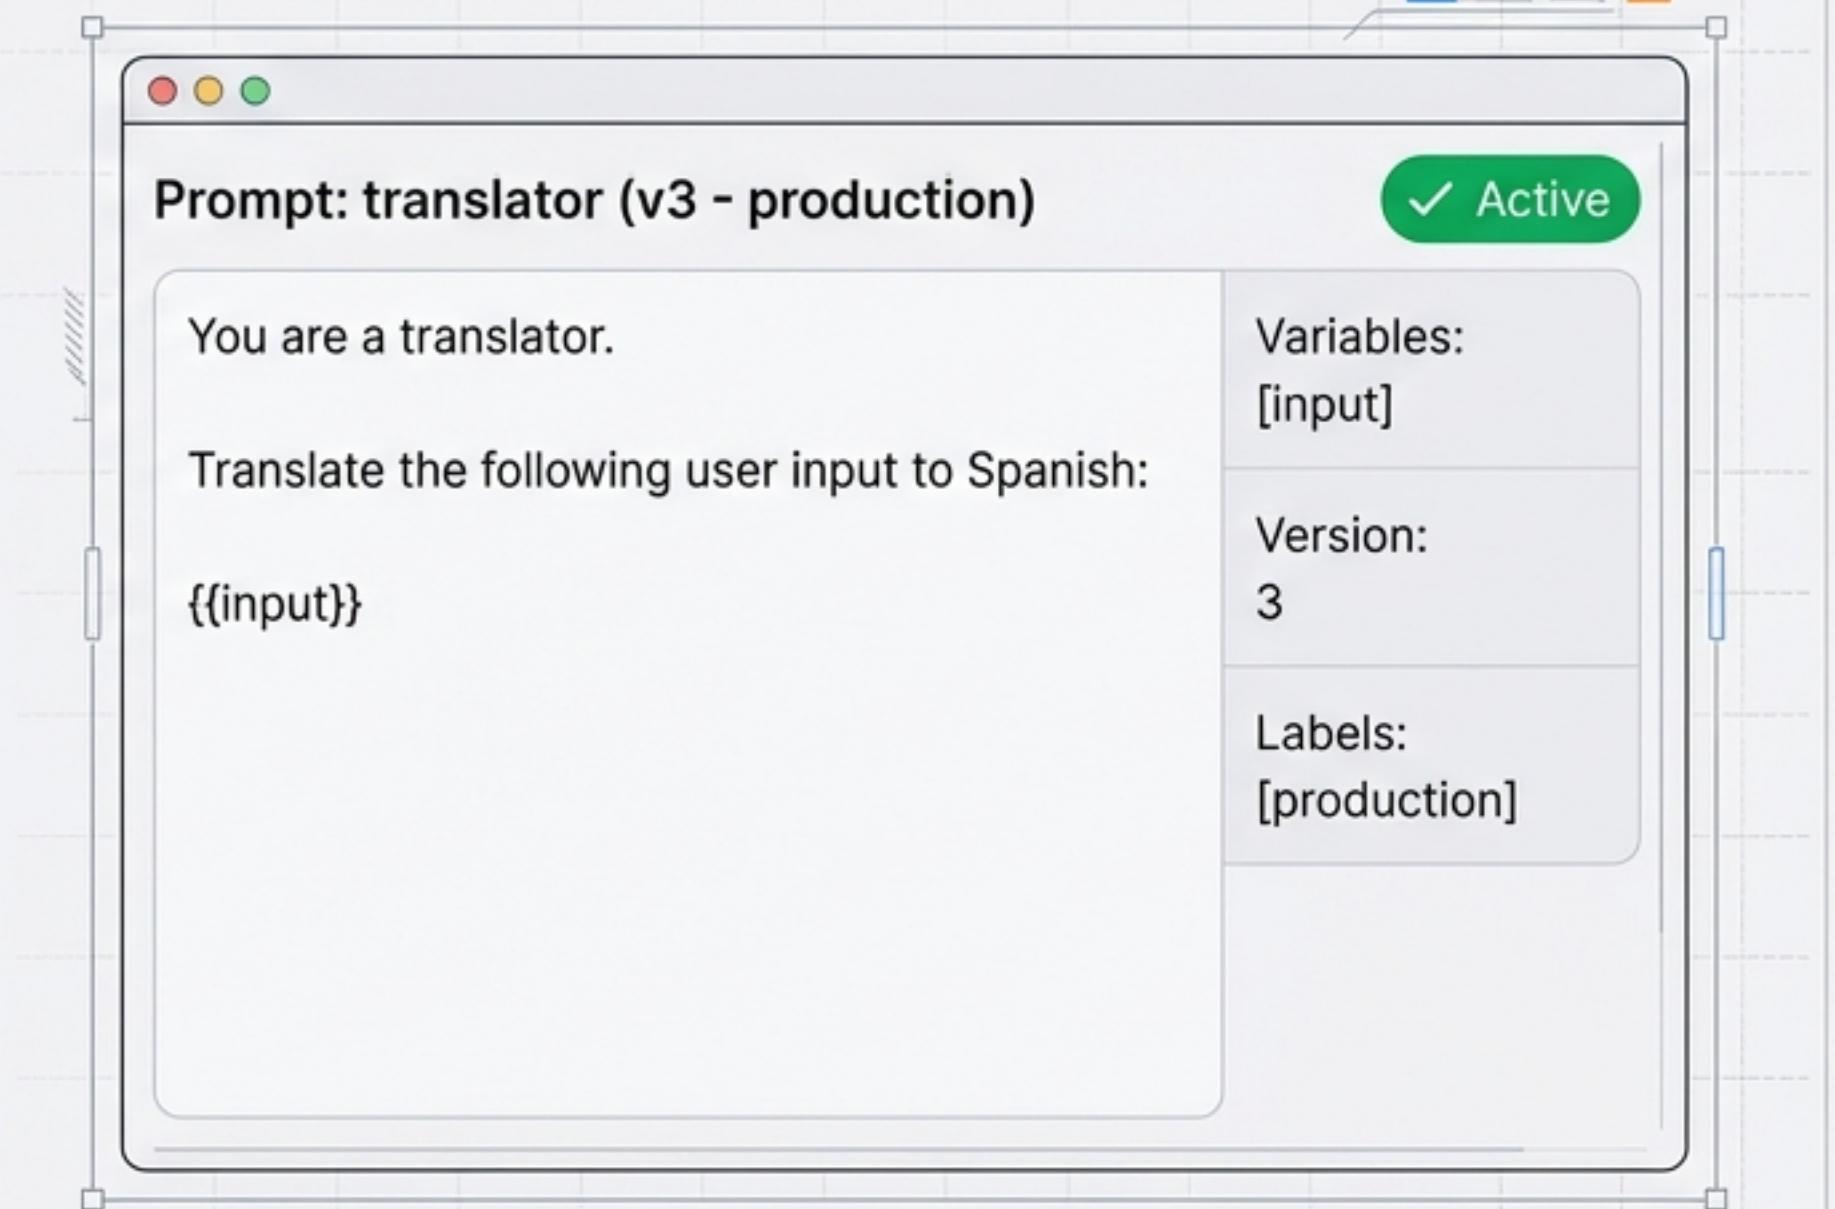
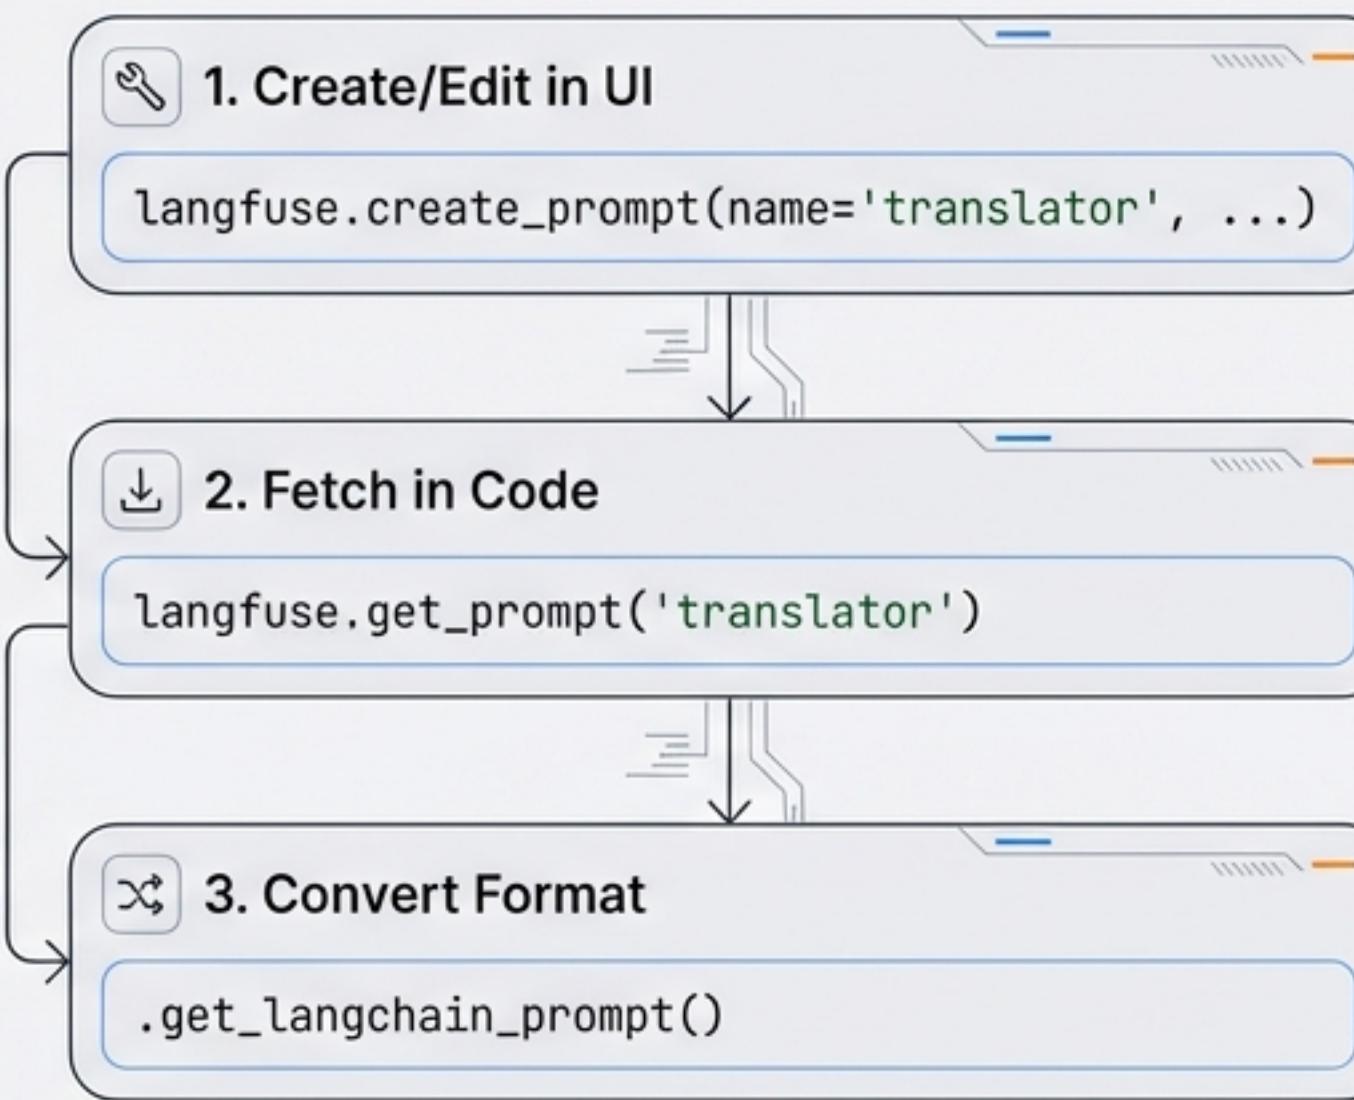
Regression testing before deployment.



Prompt Management as a CMS

Decoupling prompts from application code.

The Workflow



Prompt Editor UI Simulation

Implementation Cheatsheet

Goal

Trace a Chain

Trace a Graph Agent

Link Distributed Agents

Score a Run (Feedback)

Test a Dataset

Action / Syntax

```
chain.invoke(..., config={'callbacks':[handler]})
```

```
graph.compile().with_config({'callbacks':[handler]})
```

```
trace_context={'trace_id': predefined_id}
```

```
span.score_trace(name='feedback', value=1)
```

```
with item.run() as root_span: ...
```

Start Tracing Today

cloud.langfuse.com

```
$ pip install langfuse langchain langgraph
```

Cookbooks

Integration examples
for Python & JS

Documentation

langfuse.com/docs

GitHub

Open Source
(MIT License)

Scan for
Cookbooks

DBDI